

Reducing Write Barrier Overheads for Orthogonal Persistence



Yilin Zhang
University of Tokyo



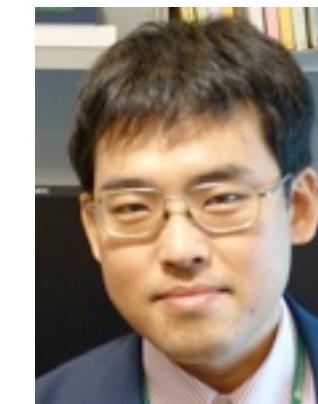
Omkar Dilip Dhawal
IIT Madras



V. Krishna Nandivada
IIT Madras



Shigeru Chiba
University of Tokyo



Tomoharu Ugawa
University of Tokyo

Reducing Write Barrier Overheads for Orthogonal Persistence

Persistence

- Non-volatile memory (NVM)
 - Intel Optane
 - Objects in byte-addressable NVM can be accessed in the same way as those in DRAM
 - Faster than SSD but slower than DRAM
 - Hybrid systems used in practice

Persistence

- Non-volatile memory (NVM)
 - Intel Optane
 - Objects in byte-addressable NVM can be accessed in the same way as those in DRAM
 - Faster than SSD but slower than DRAM
 - Hybrid systems used in practice
- Issues
 - Which objects must be made persistent ? When should the object become persistent ?
 - Tedious and error-prone task for programmers

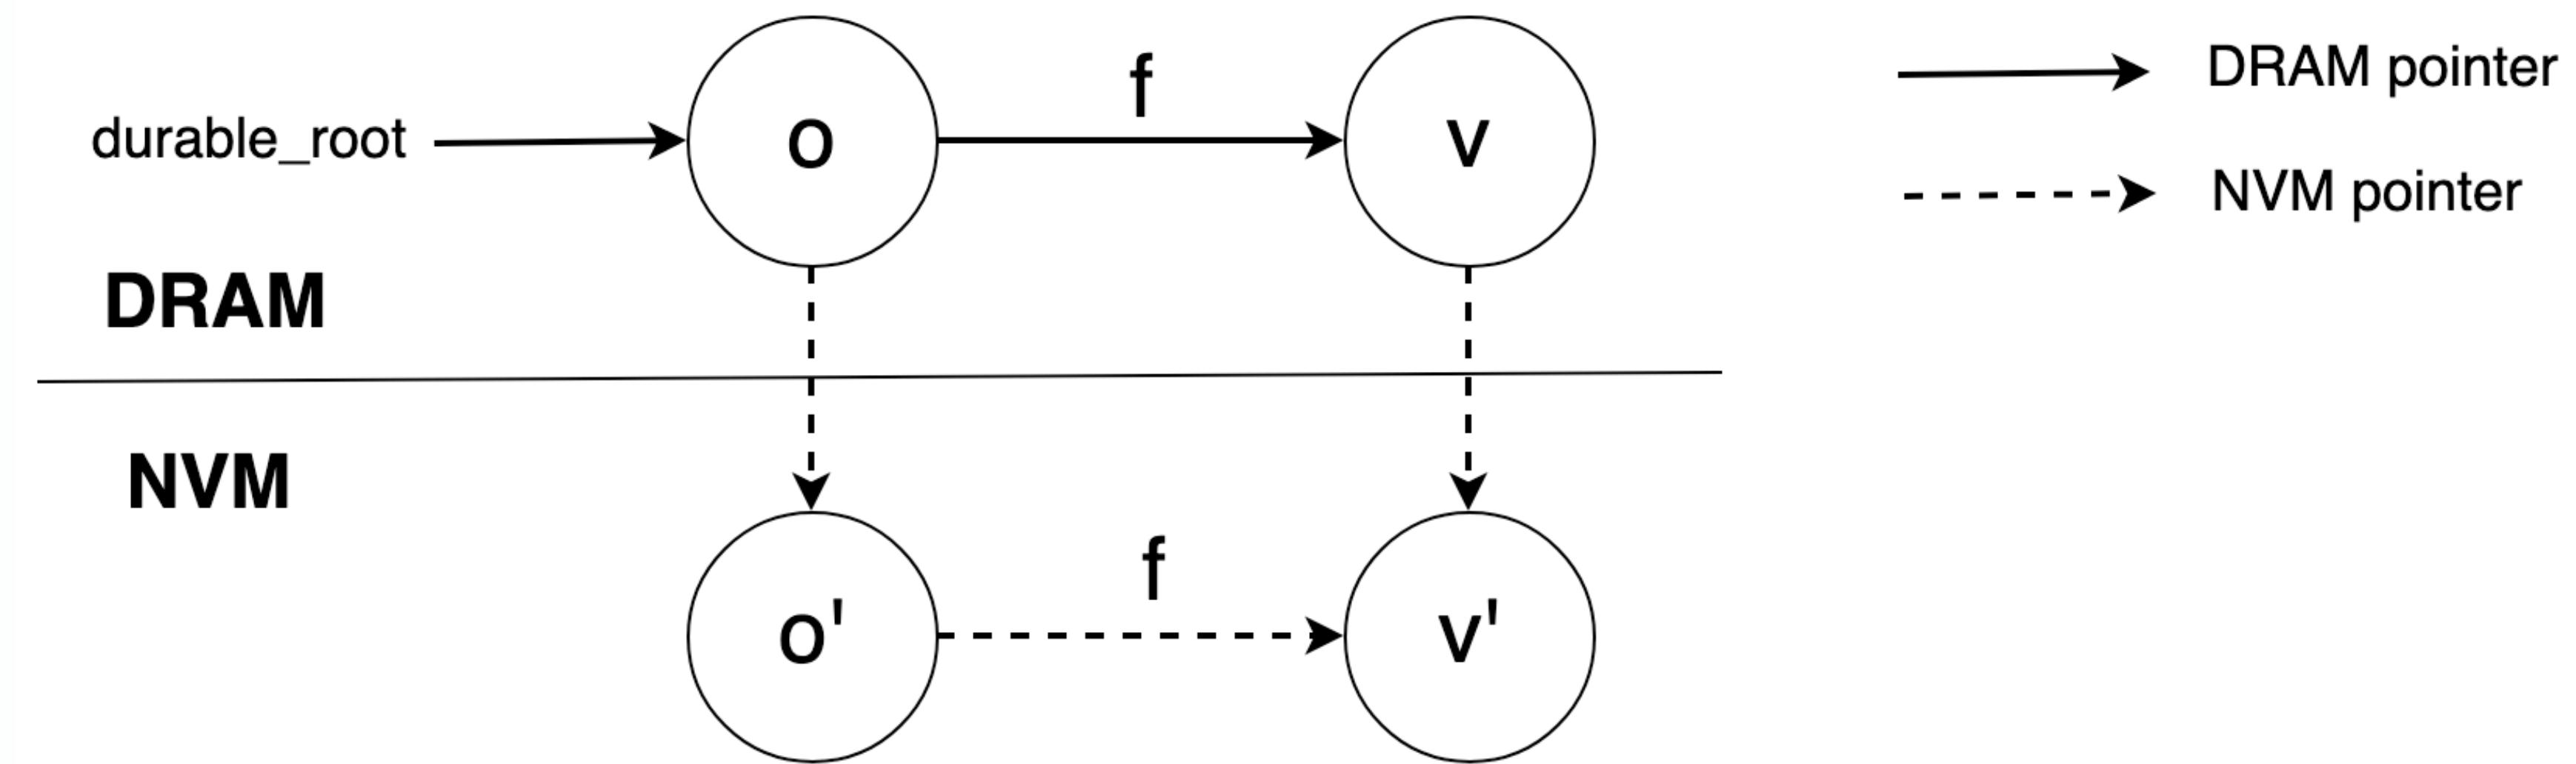
Reducing Write Barrier Overheads for Orthogonal Persistence

Orthogonal Persistence

- Programmers can annotate static fields as durable roots
- Persistence of objects decided by reachability

Orthogonal Persistence

- Programmers can annotate static fields as durable roots
- Persistence of objects decided by reachability
- Objects reachable from persistent roots are copied to NVM without programmer's intervention (Replication-based object persistence) [Matsumoto et al. 2022]



Orthogonal Persistence

- Persistence of objects decided by reachability
- Programmers can annotate static fields as persistent roots
- Objects reachable from persistent roots are copied to NVM without programmer's intervention
- Issues
 - Java supports multi-threading
 - Concurrent access: One thread is modifying an object while another thread is attempting to copy it to NVM.

Concurrent Access

Copier Thread

durable_root = o

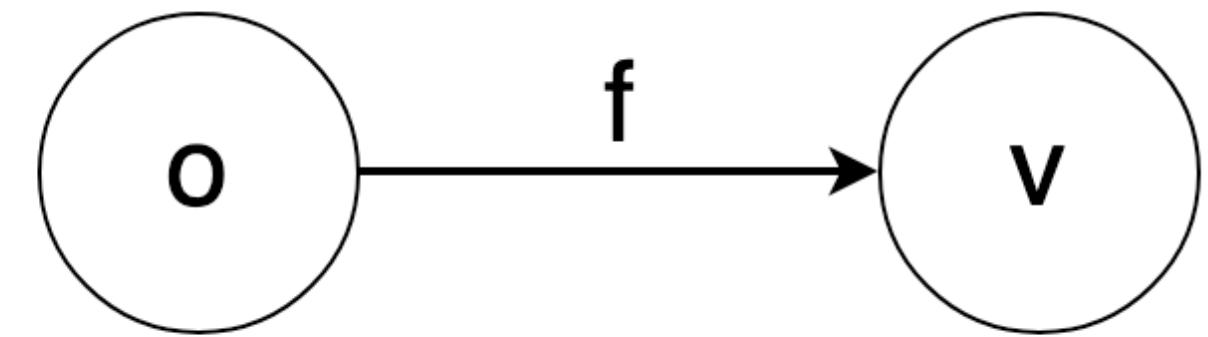
Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 1: Replica of o is absent

DRAM

NVM



Concurrent Access

Copier Thread

durable_root = o

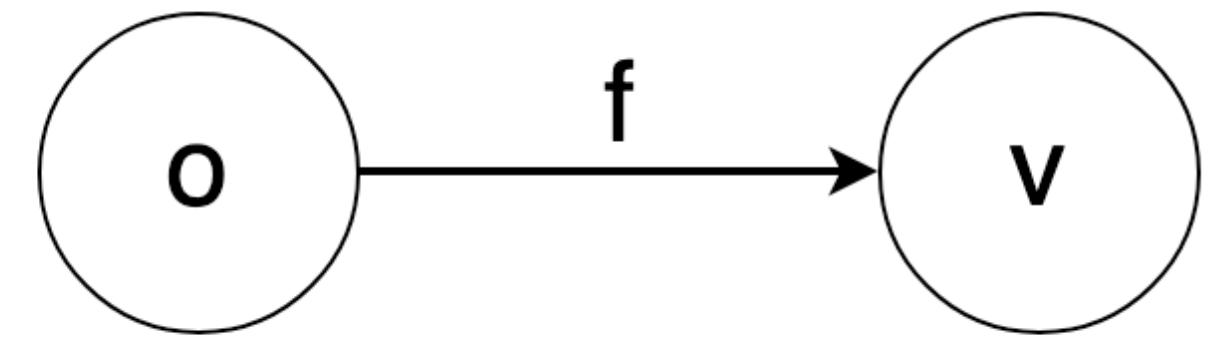
Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 1: Replica of o is absent

DRAM

NVM



Concurrent Access

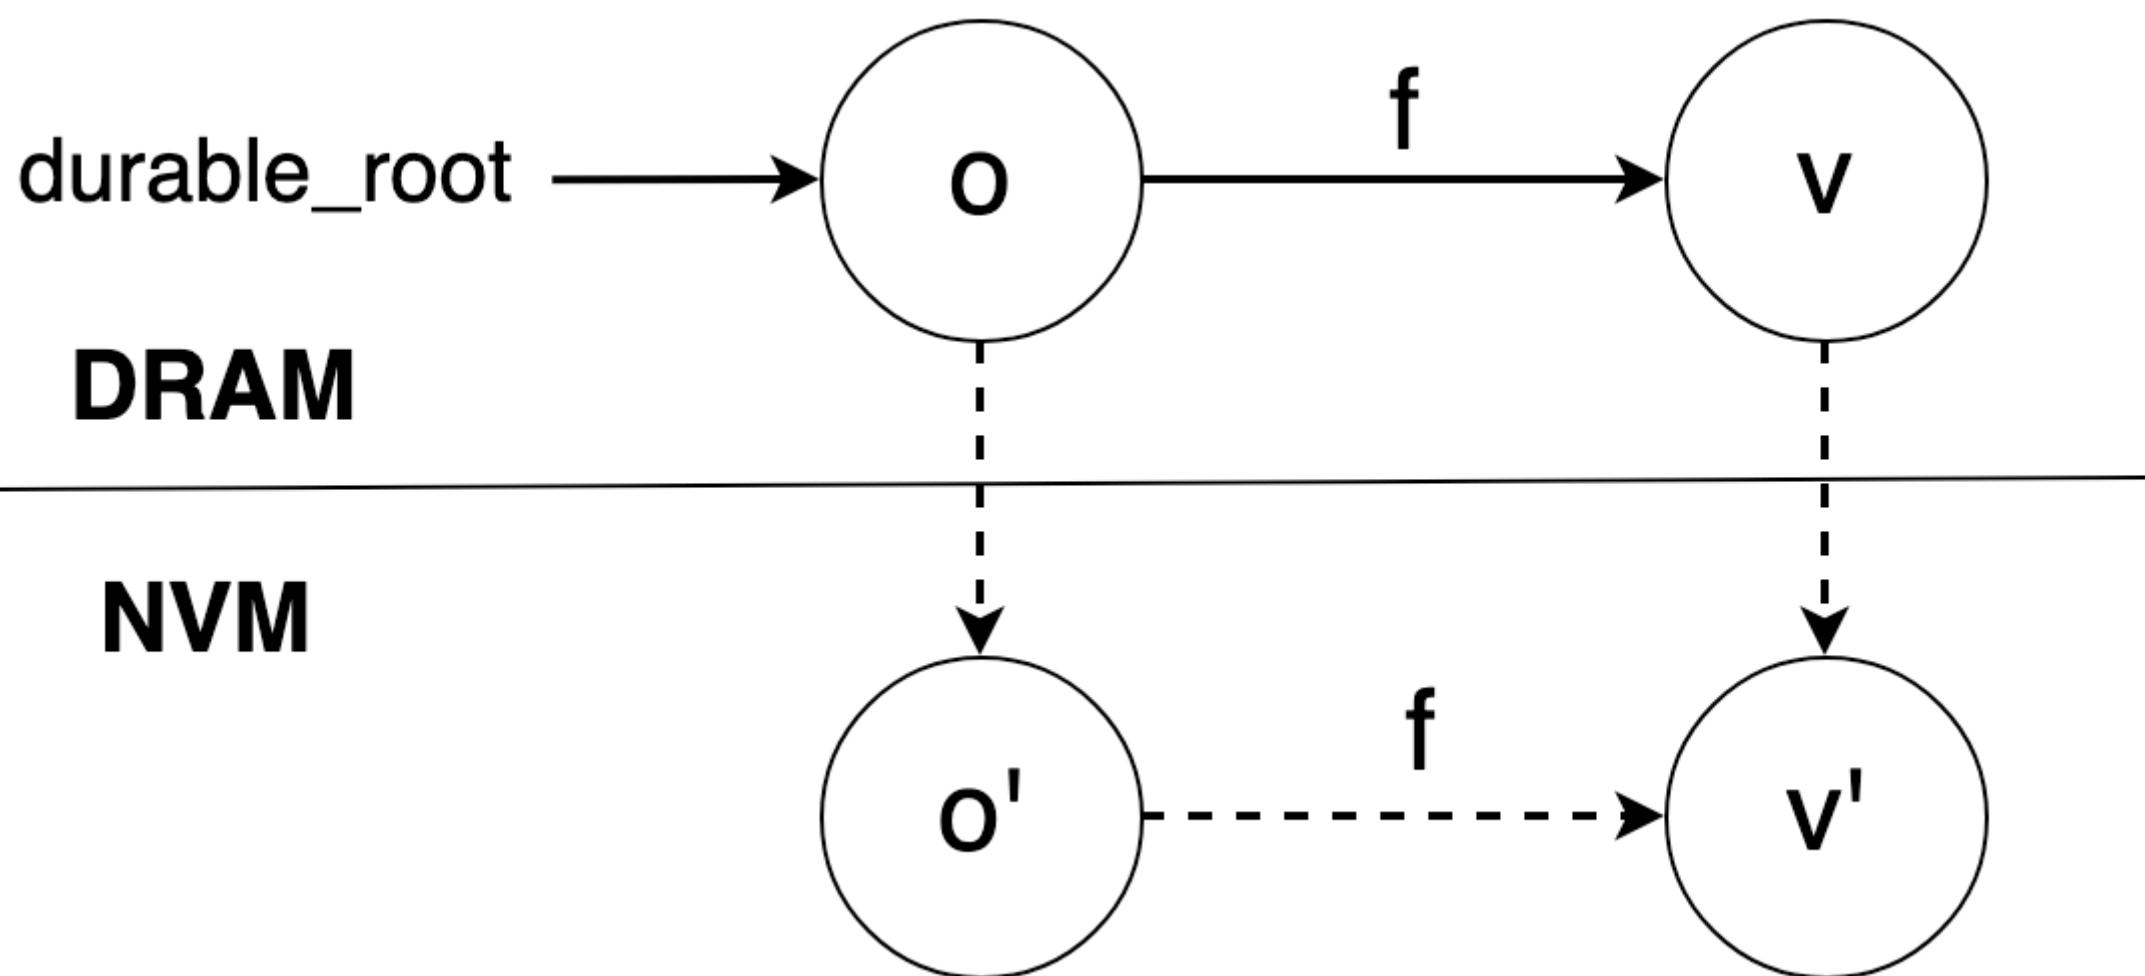
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 1: Replica of o is absent



Concurrent Access

Case 2: Replica of o is already present

Concurrent Access

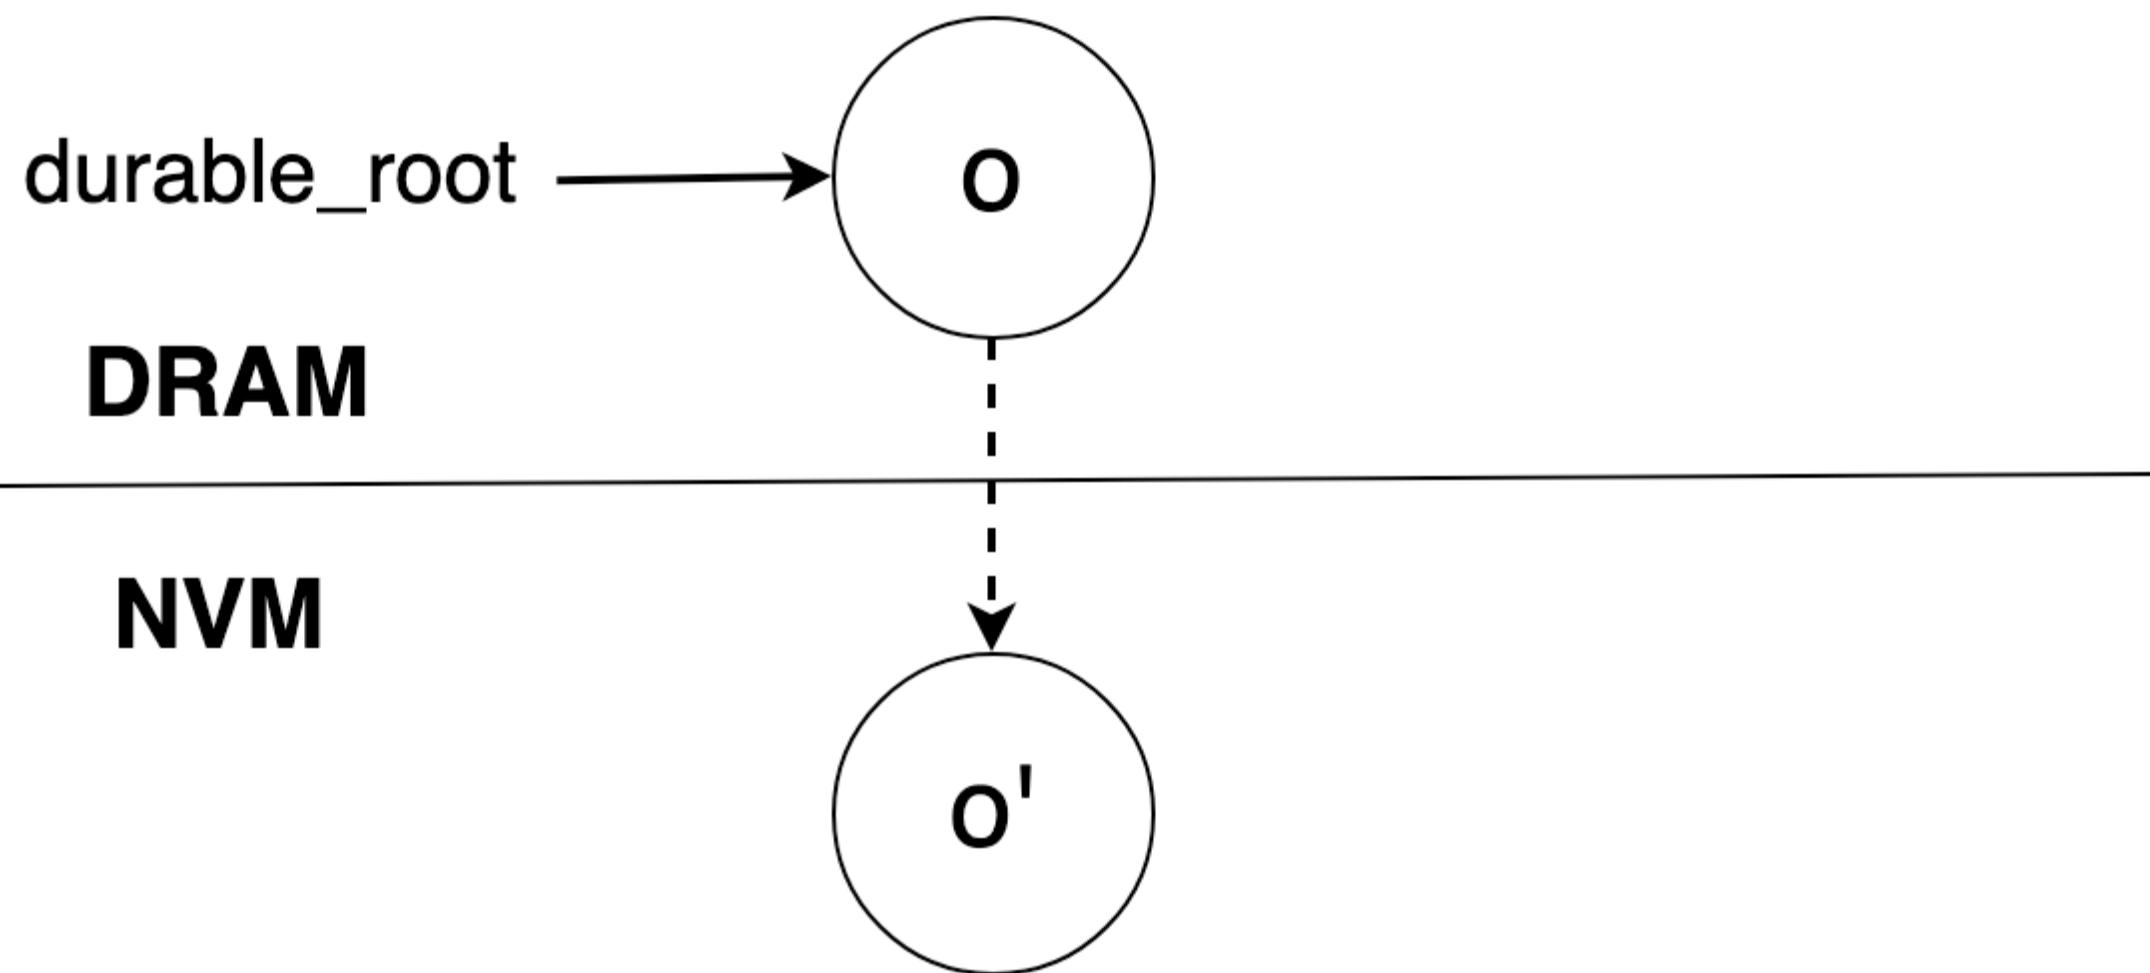
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 2: Replica of o is present



Concurrent Access

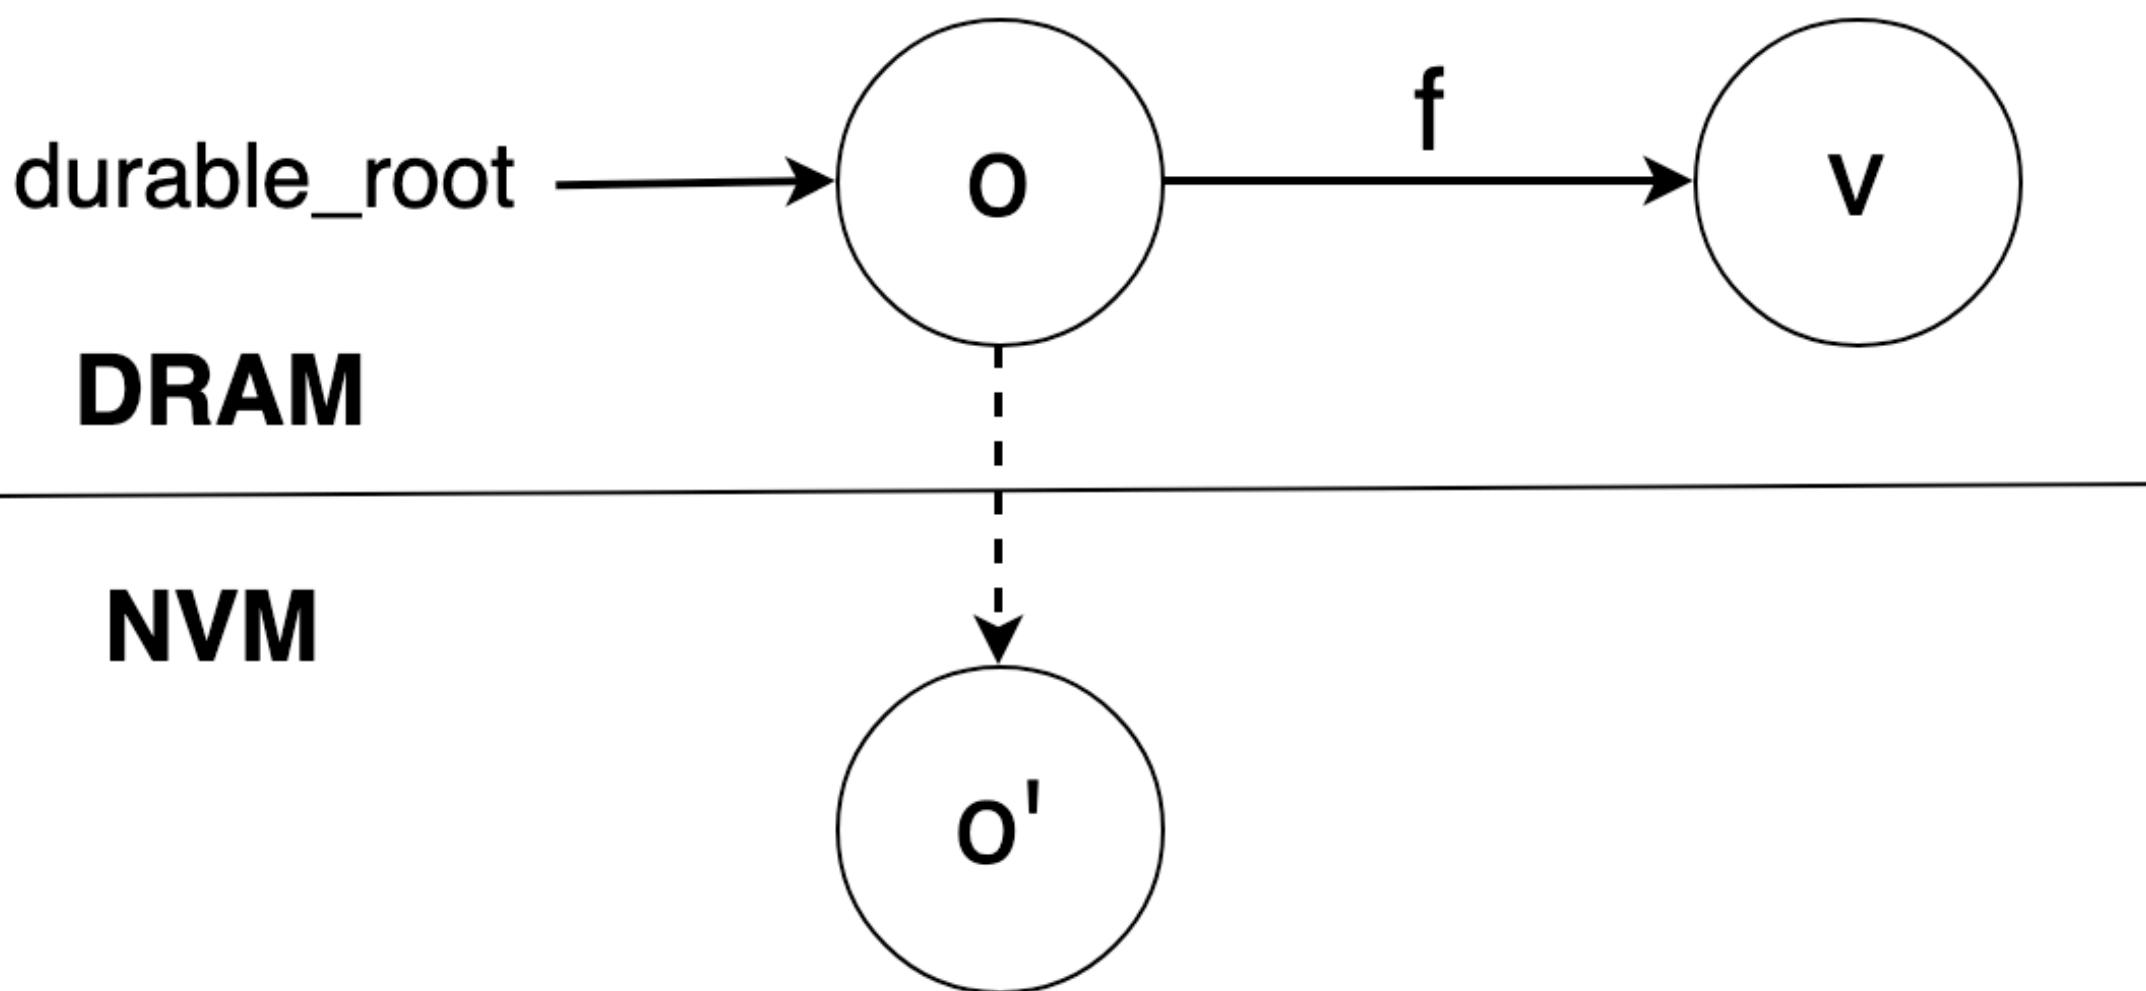
Copier Thread

durable_root = o

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 2: Replica of o is present



Concurrent Access

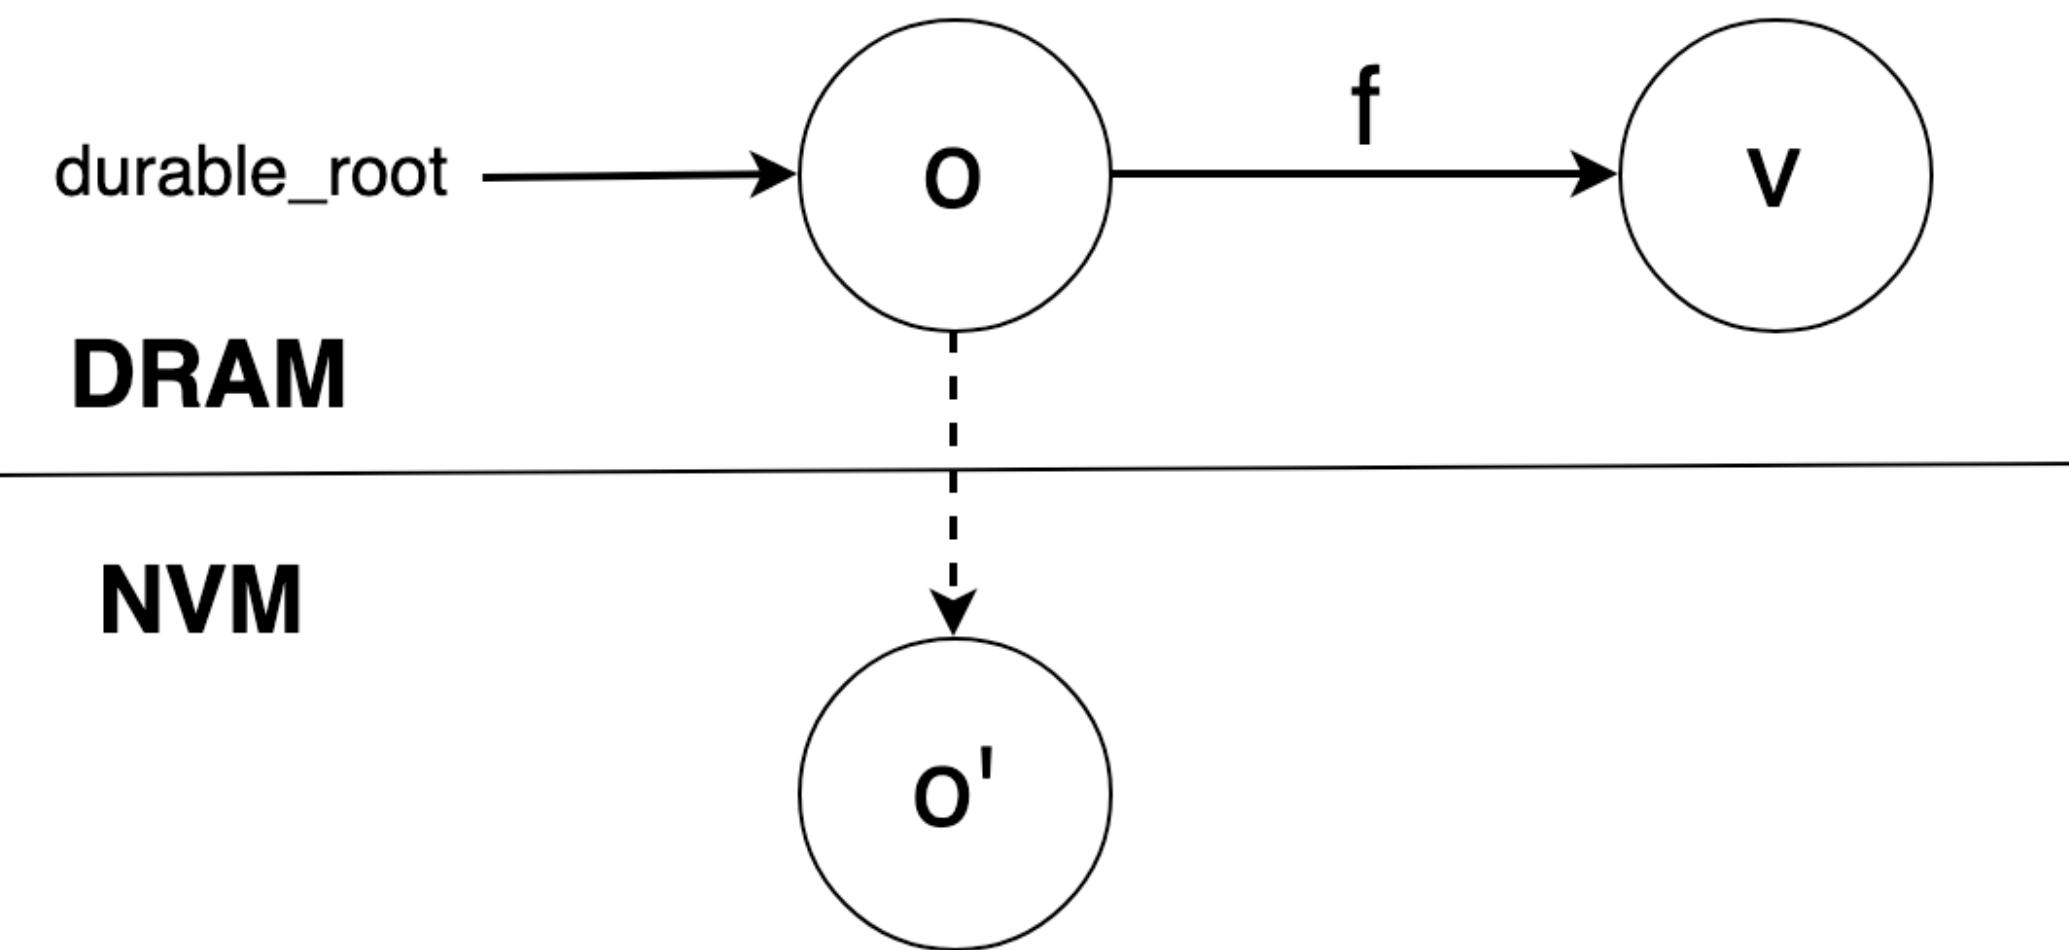
Copier Thread

durable_root = o

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 2: Replica of o is present



Concurrent Access

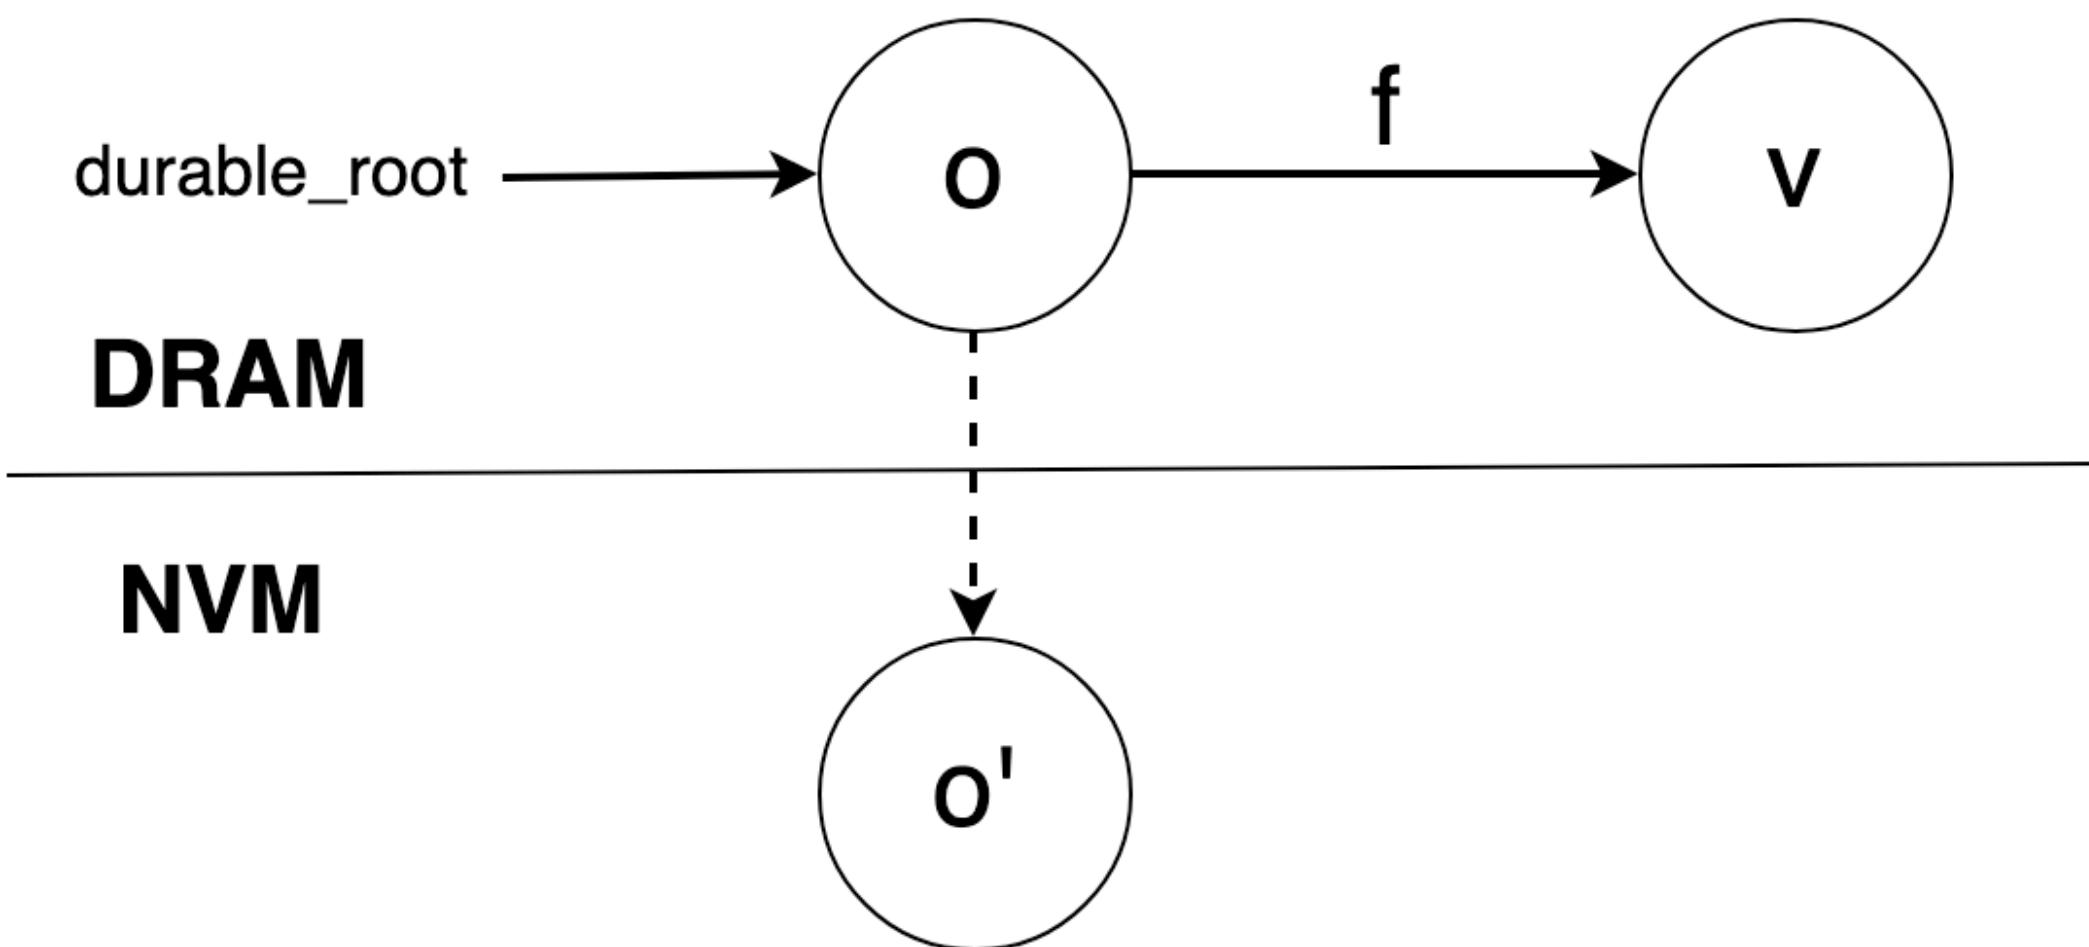
Copier Thread

durable_root = o

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 2: Replica of o is present



Concurrent Access

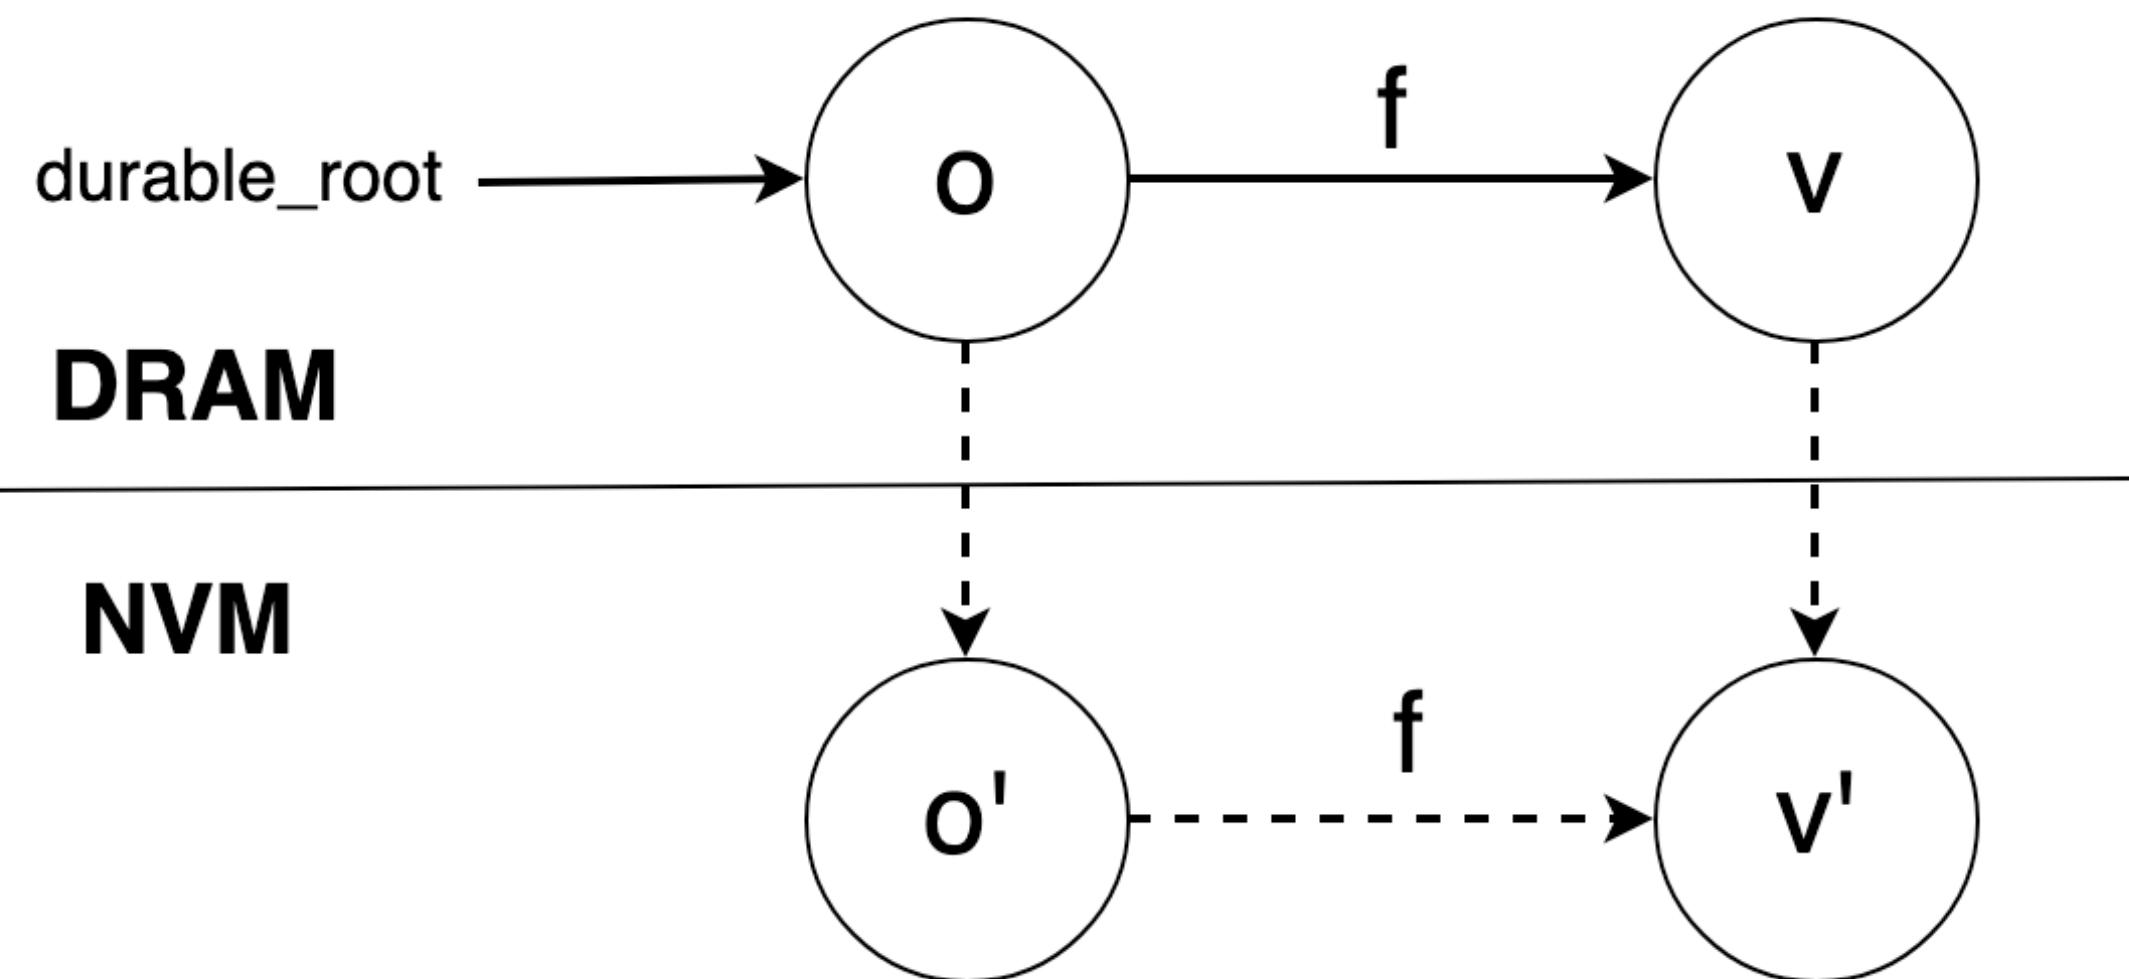
Copier Thread

durable_root = o

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 2: Replica of o is present



Concurrent Access

**Case 3: Instructions are reordered
(x86 Weak memory consistency model)**

Concurrent Access

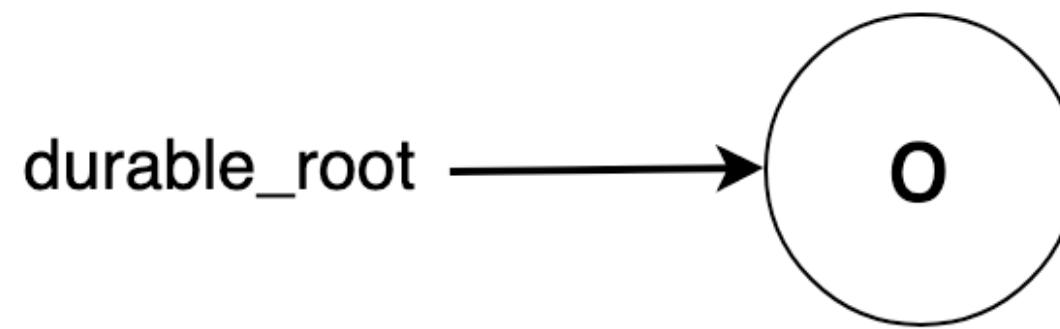
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



NVM

Concurrent Access

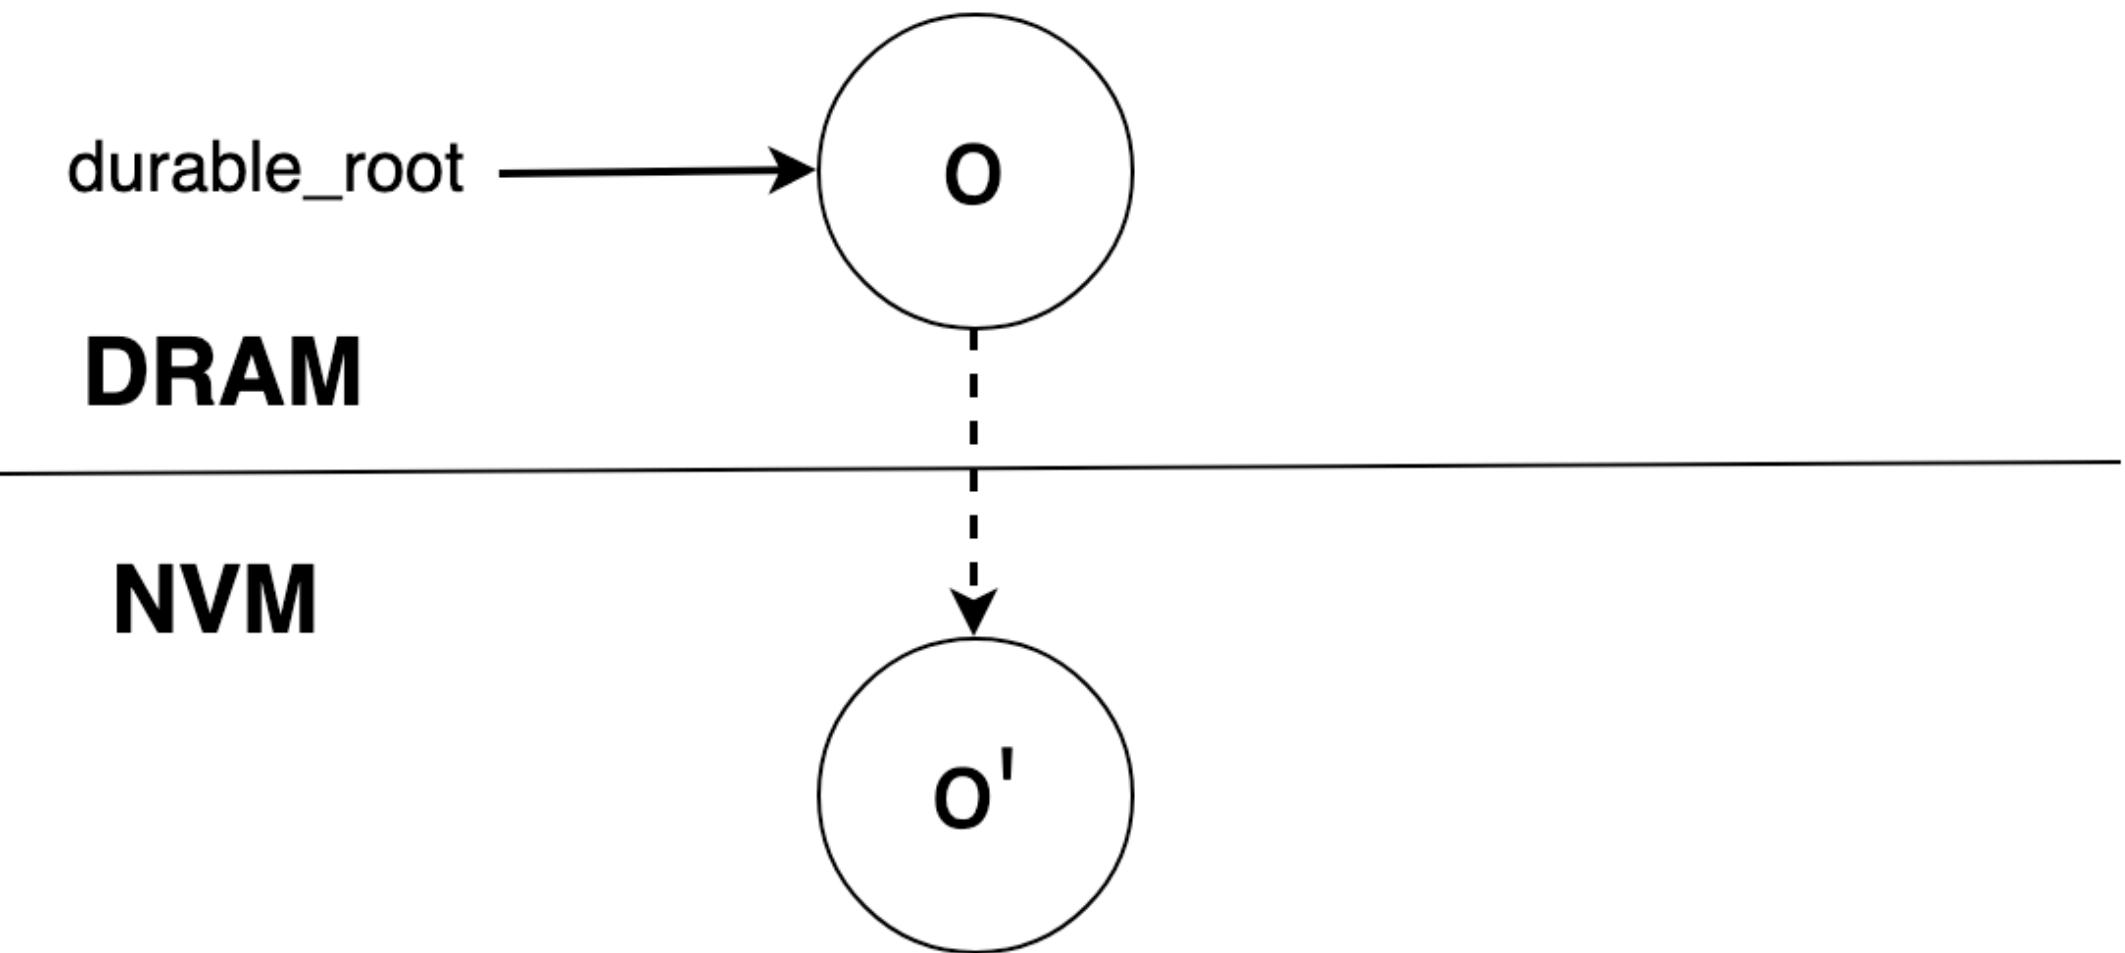
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Concurrent Access

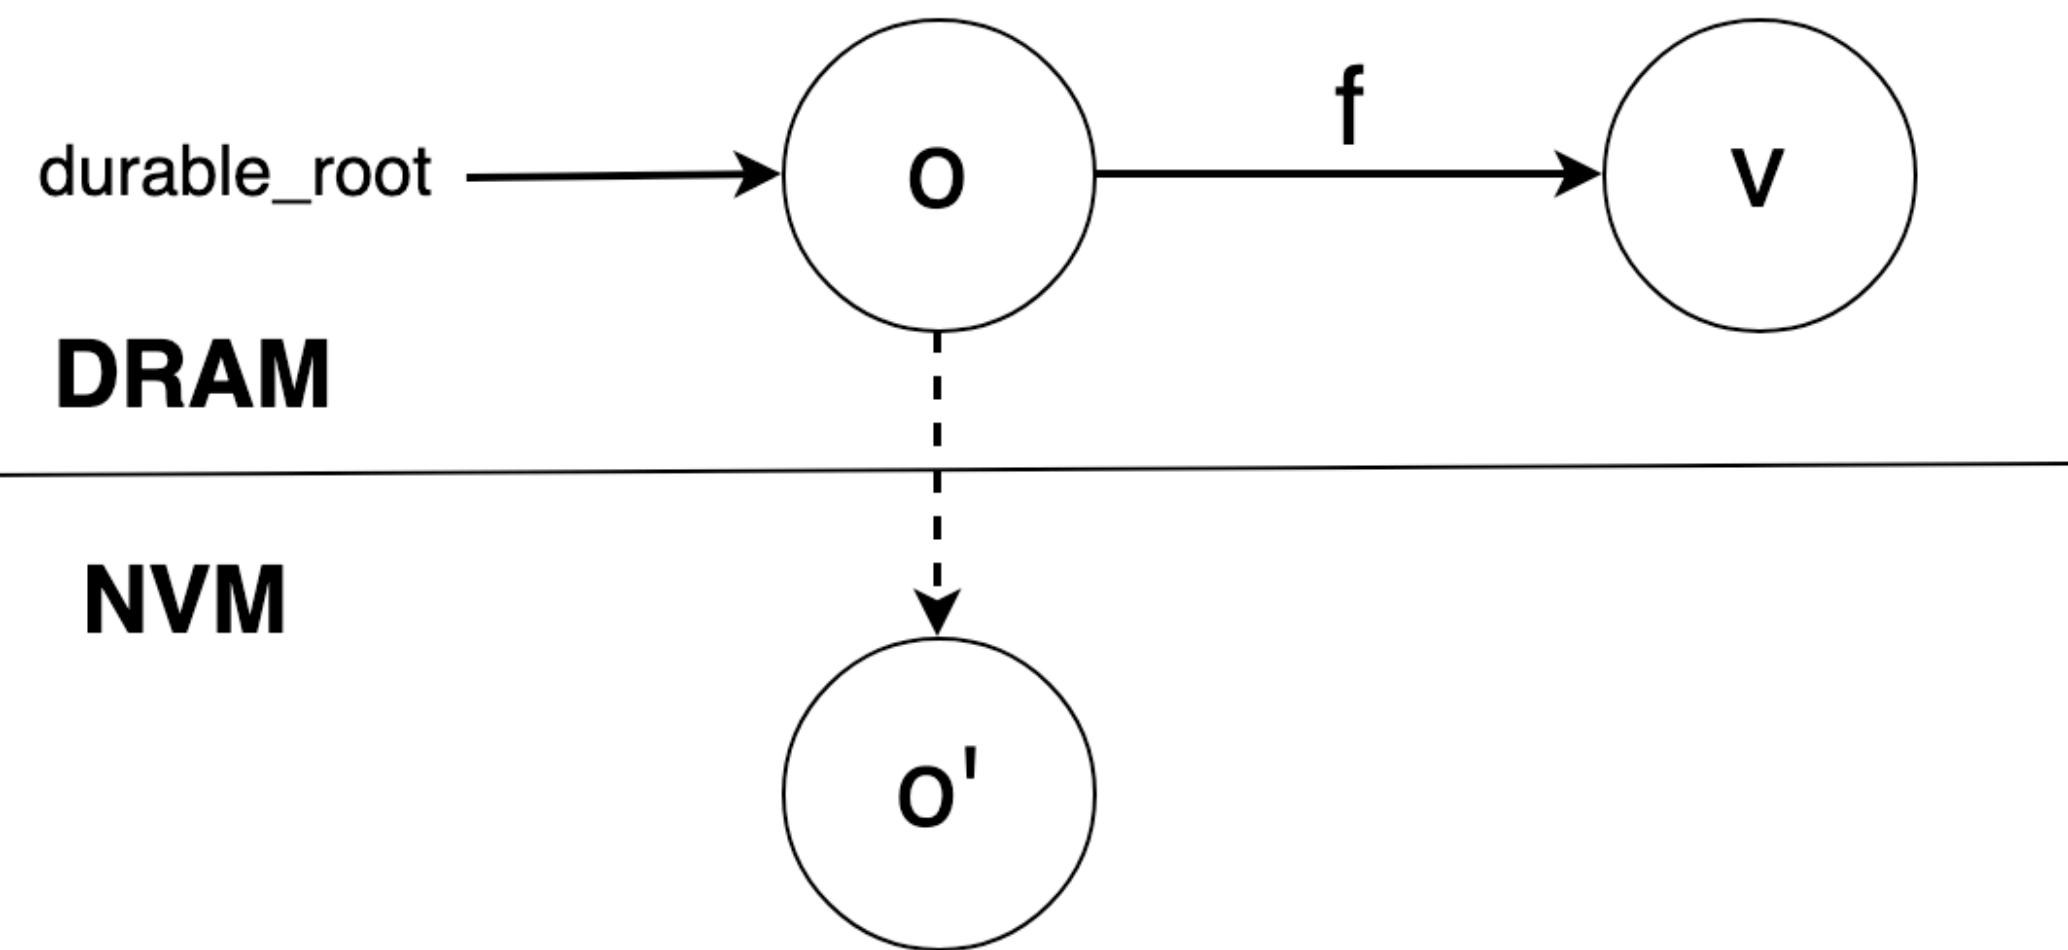
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Concurrent Access

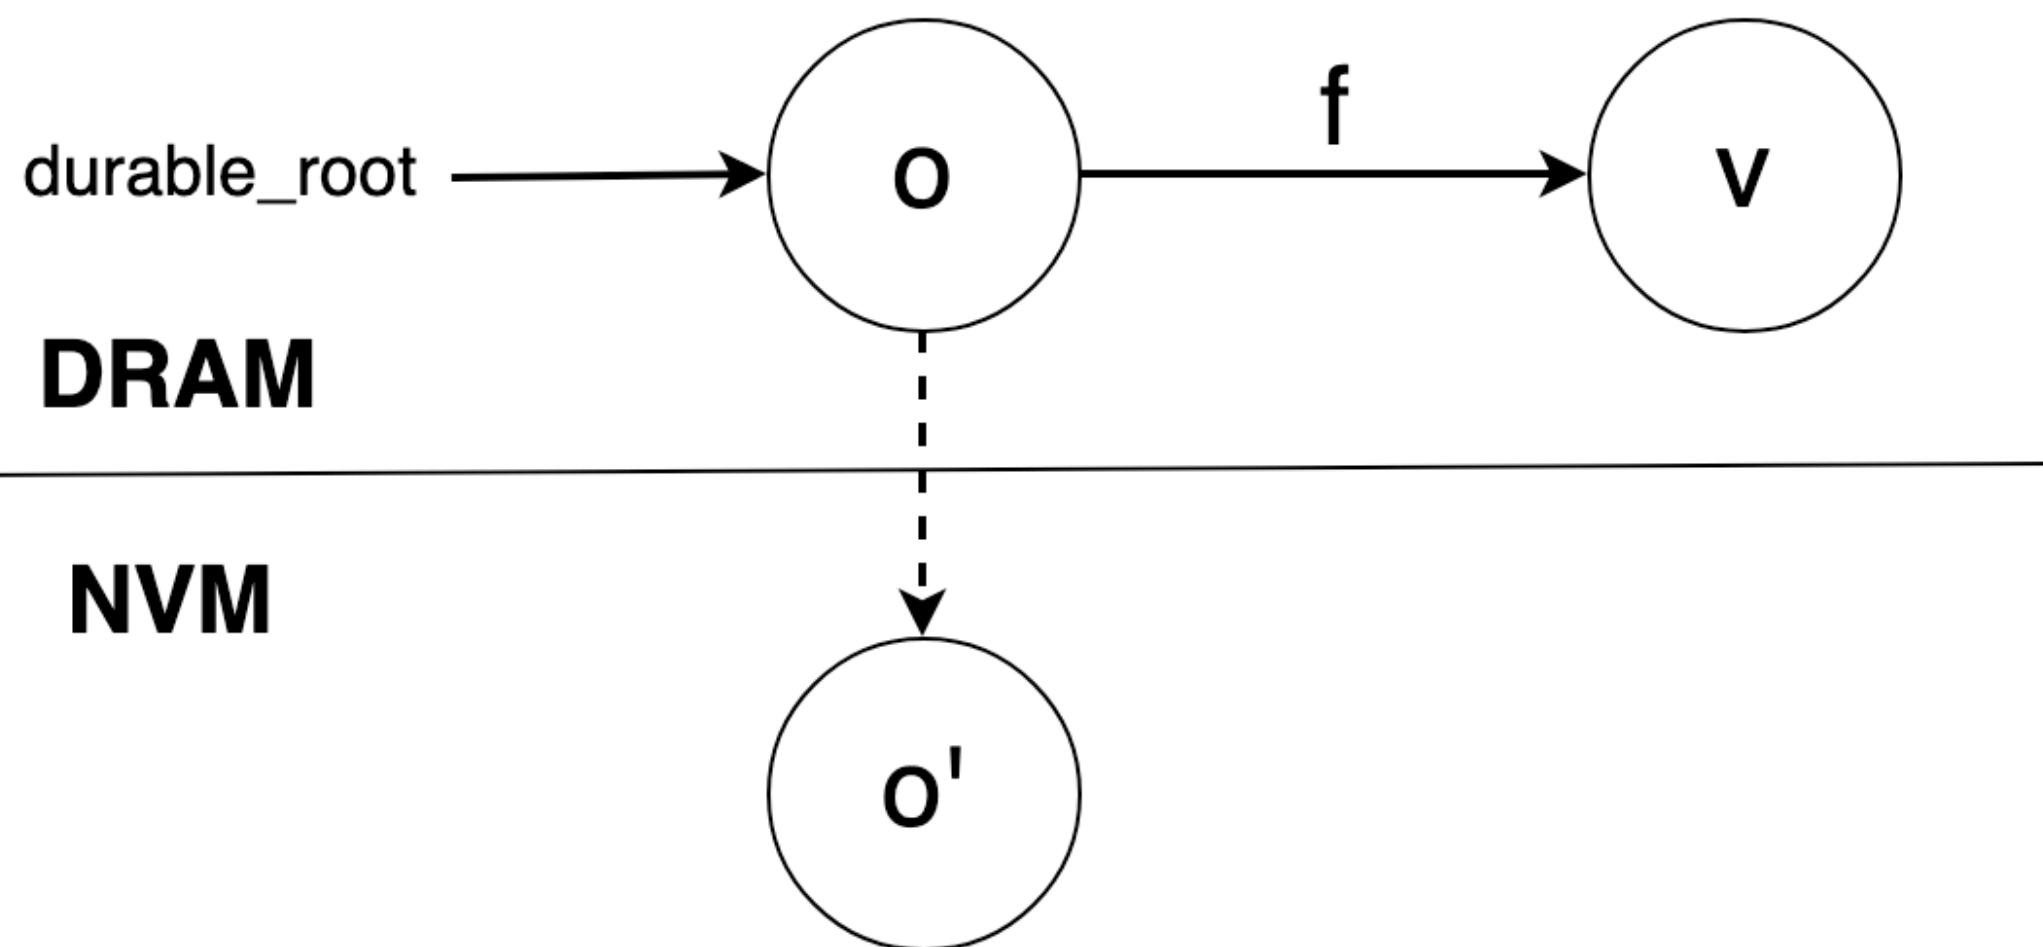
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Concurrent Access

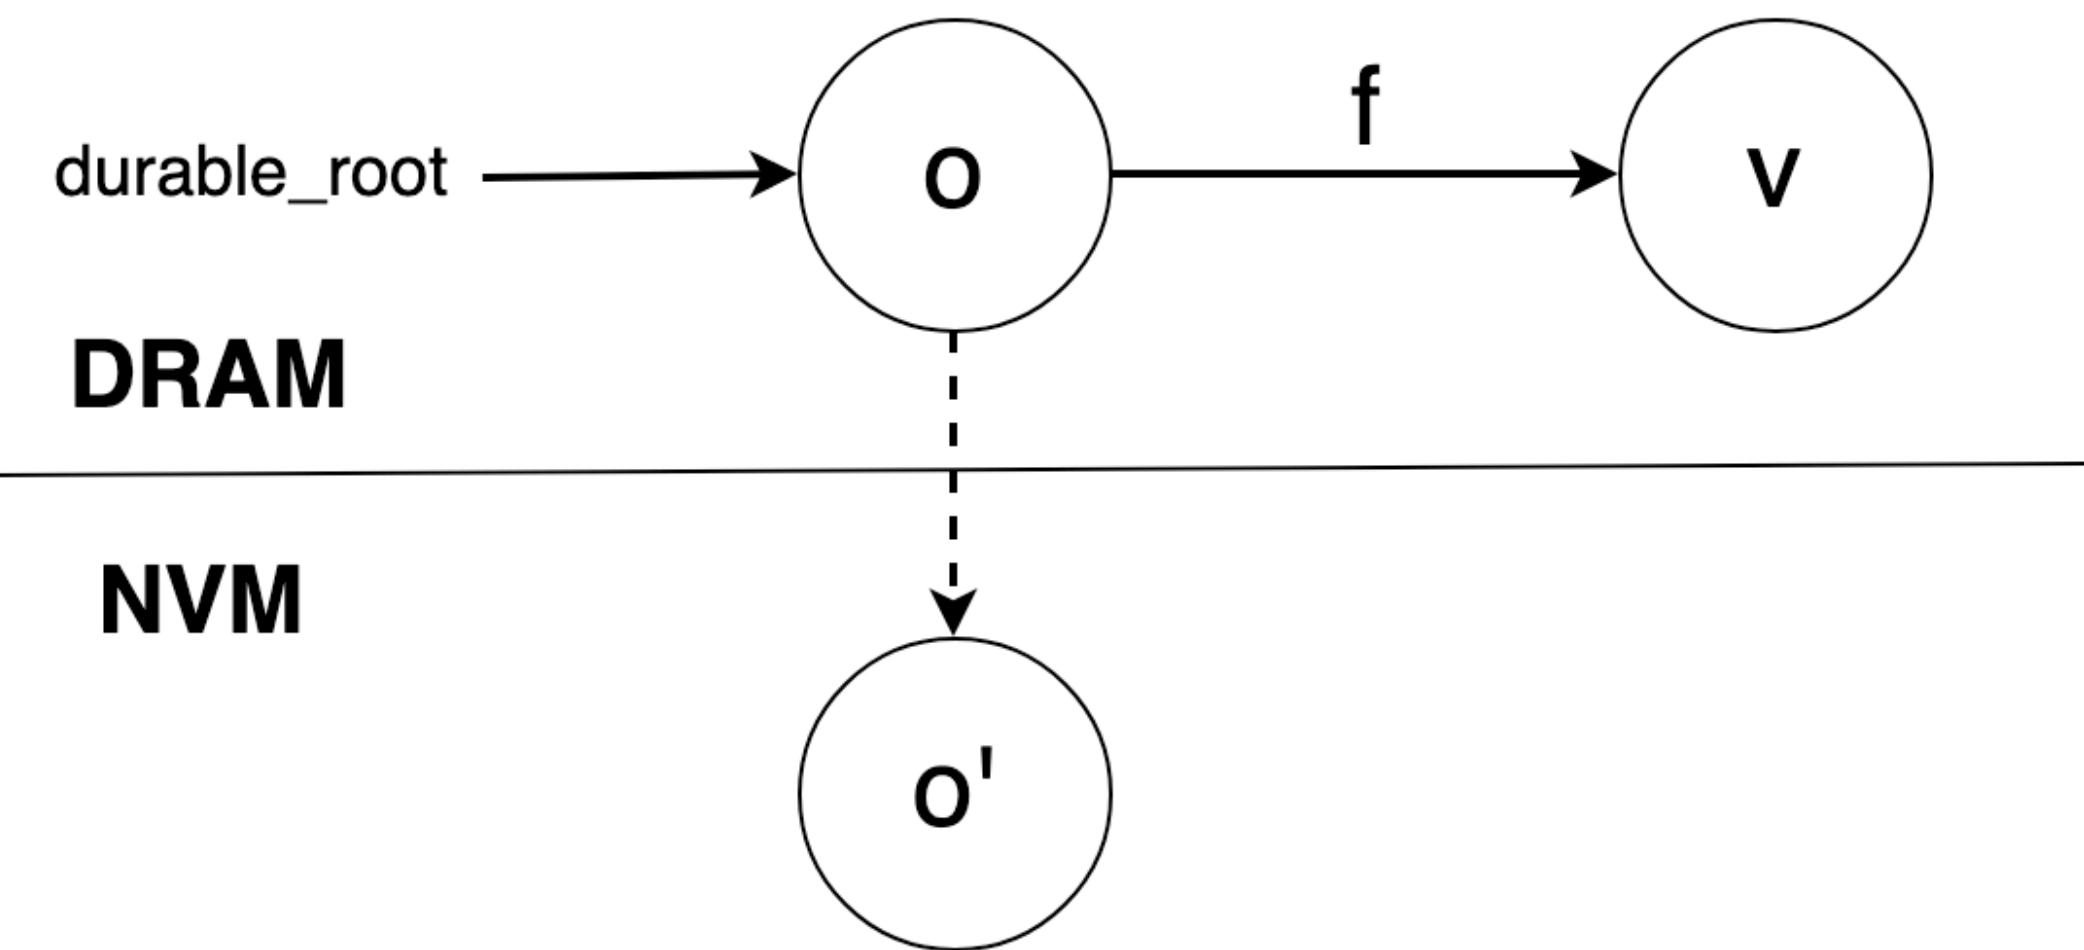
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Copies inconsistent !!

[Matsumoto et al. 2022]

Reducing Write Barrier Overheads for Orthogonal Persistence

Write Barrier

Copier Thread

durable_root = o

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3     MFENCE  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

- **MFENCE** ensures all preceding load and store instructions become globally visible before any that follow it.
- Writer-wait approach
- **Issue** - MFENCE is executed for all field write instructions.

Write Barrier

Copier Thread

durable_root = o

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3     MFENCE  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

- **MFENCE** ensures all preceding load and store instructions become globally visible before any that follow it.
- Writer-wait approach

Write Barrier

Copier Thread

durable_root = o

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3     MFENCE  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

- **MFENCE** ensures all preceding load and store instructions become globally visible before any that follow it.
- Writer-wait approach
- **Issue** - MFENCE is executed for all field write instructions.

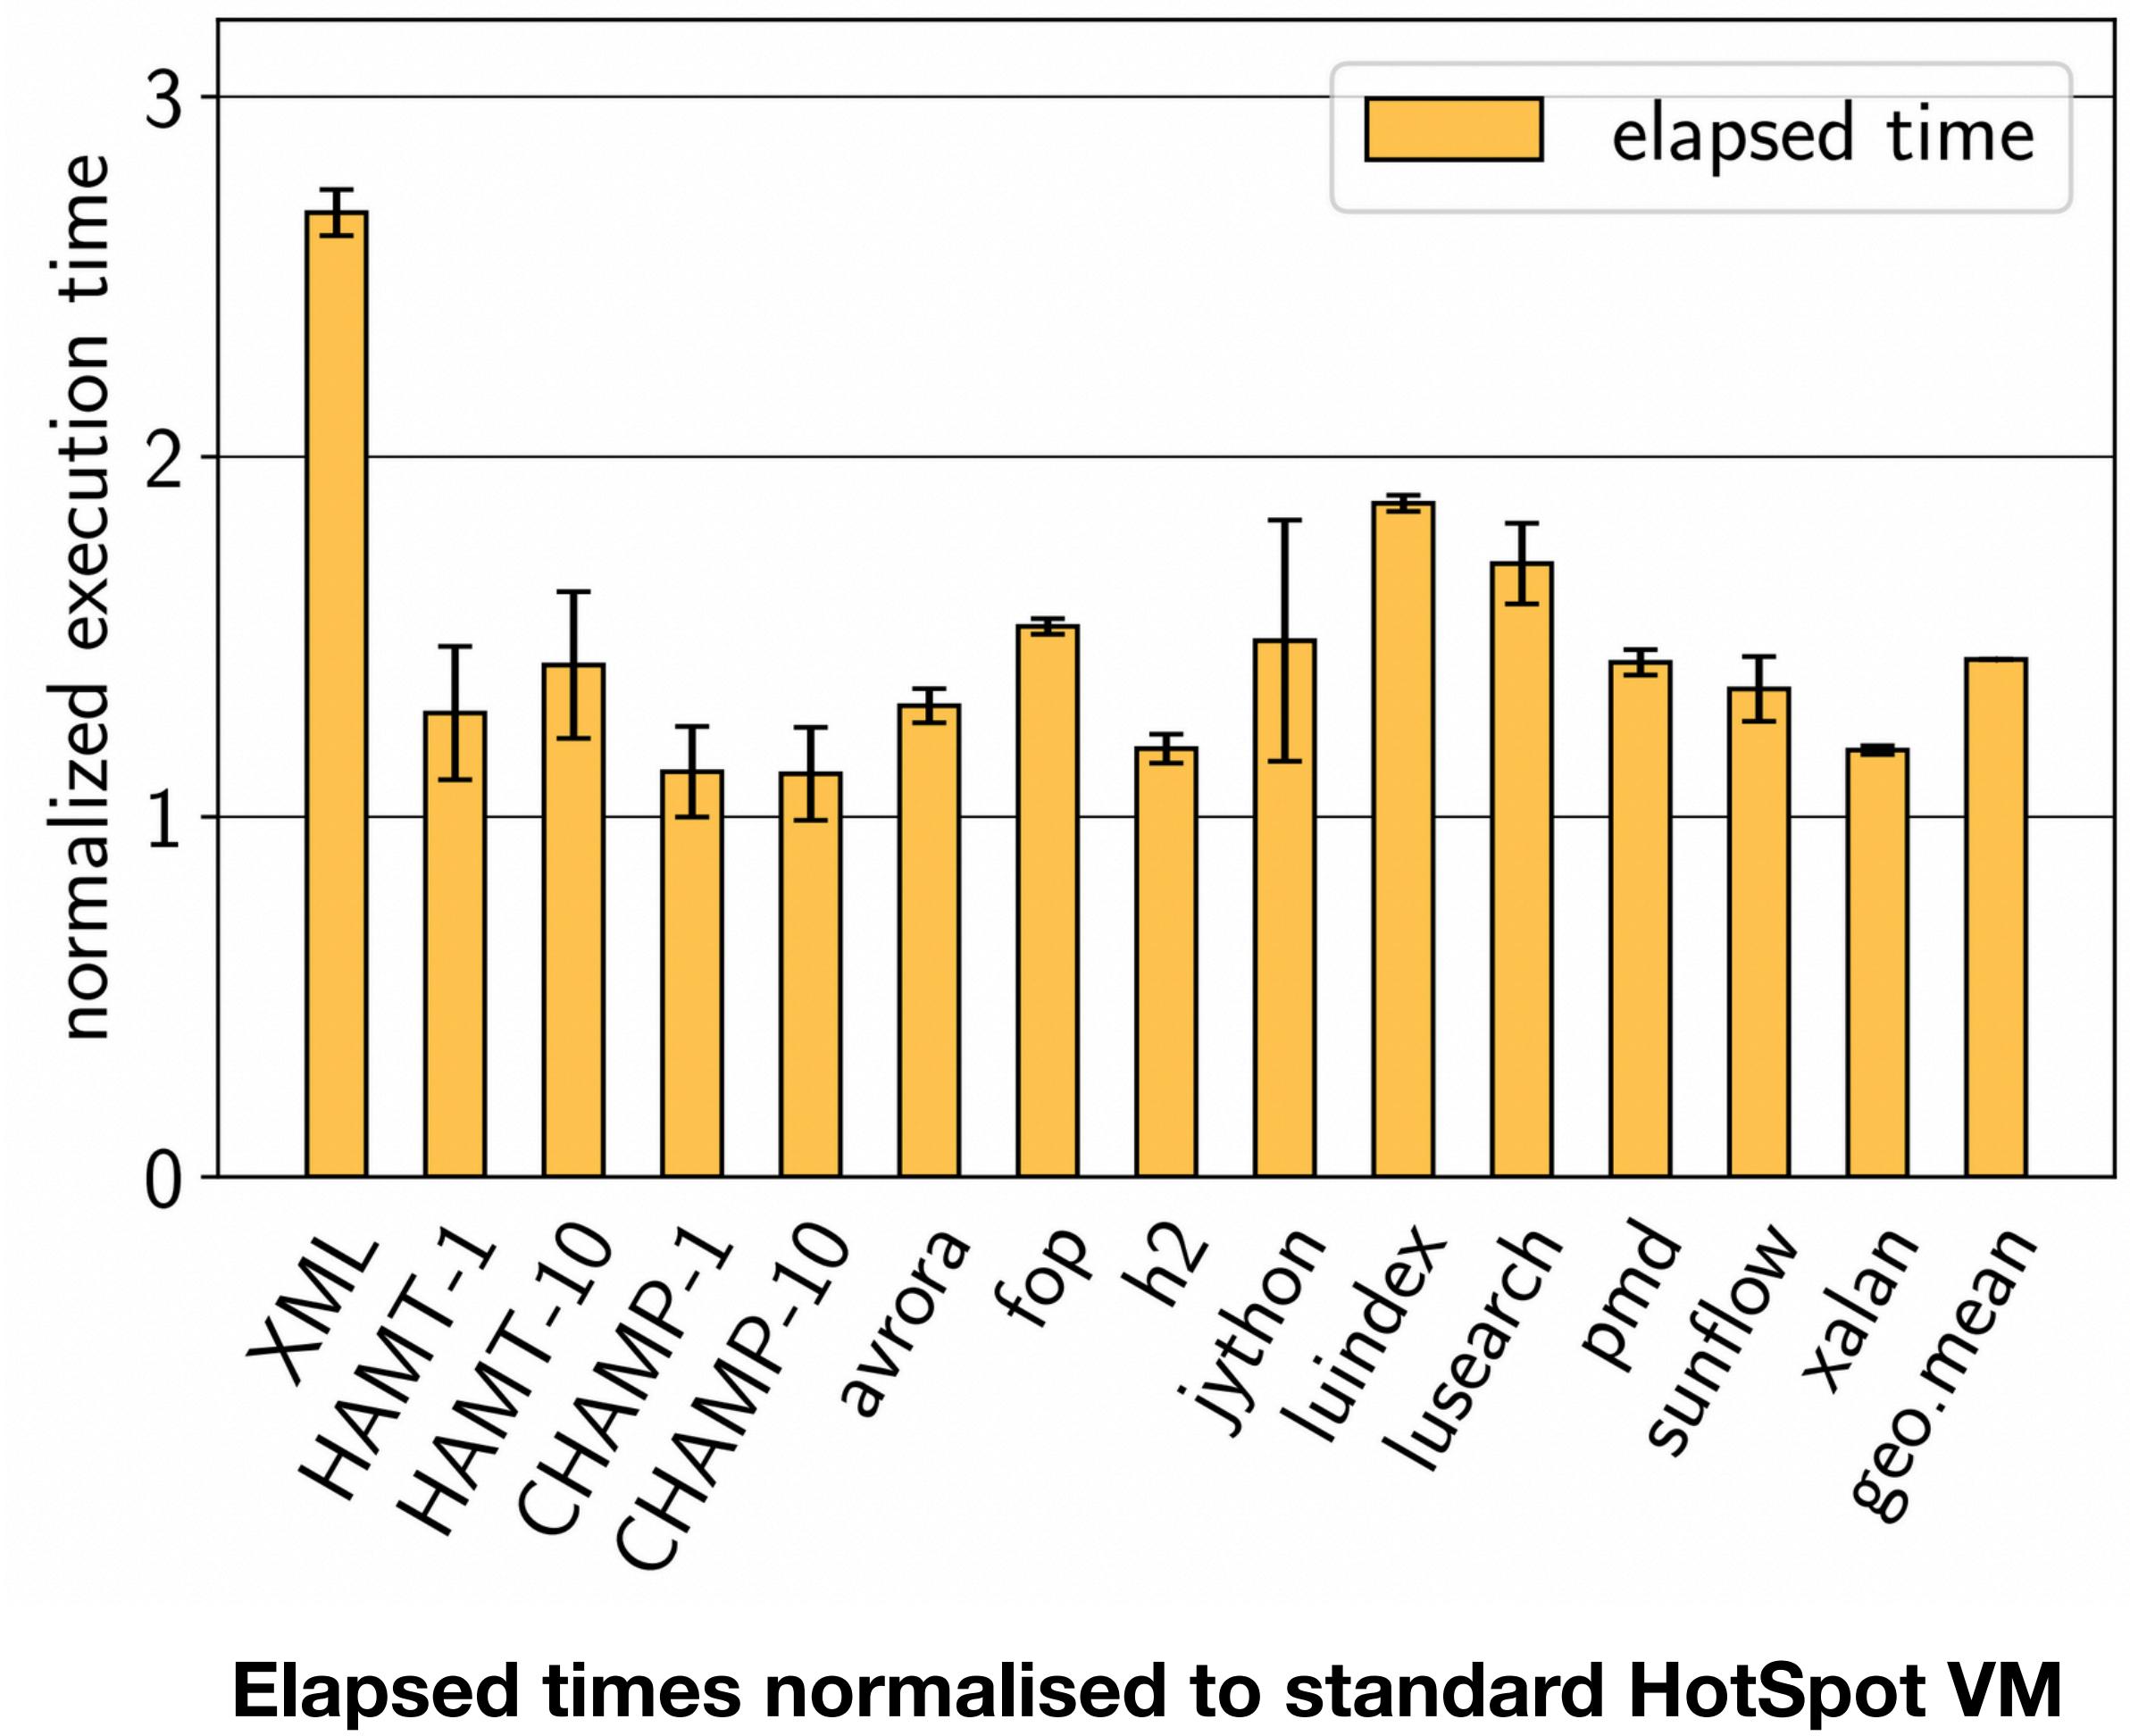
Reducing Write Barrier Overheads for Orthogonal Persistence

Write Barrier Overhead

- All putfield instructions will execute the write barrier
- Even if durable roots are absent, write barrier is executed.
- Increases execution time of the program

Write Barrier Overhead

- All putfield instructions will execute the write barrier
- Even if durable roots are absent, write barrier is executed.
- Increases execution time of the program
- Average overhead of 43.7% on the benchmarks in the absence of durable roots.



Reducing Write Barrier Overheads for Orthogonal Persistence

Intuition

- Frequency of copying << Frequency of writing
 - Shift the overhead to copier
- Copier thread performs a **Handshake** with all the threads and waits for acknowledgement
- Copier thread performs copy only when **Handshake** is acknowledged

Concurrent Access

Case 3: Read at Line 4 is reordered with Write at Line 2

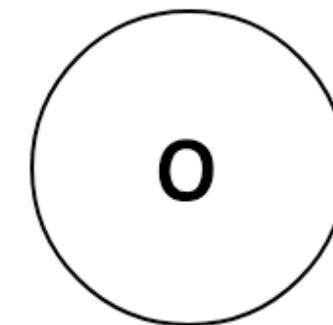
Copier Thread

```
durable_root = o
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

DRAM



NVM

Concurrent Access

Copier Thread

```
durable_root = o
```

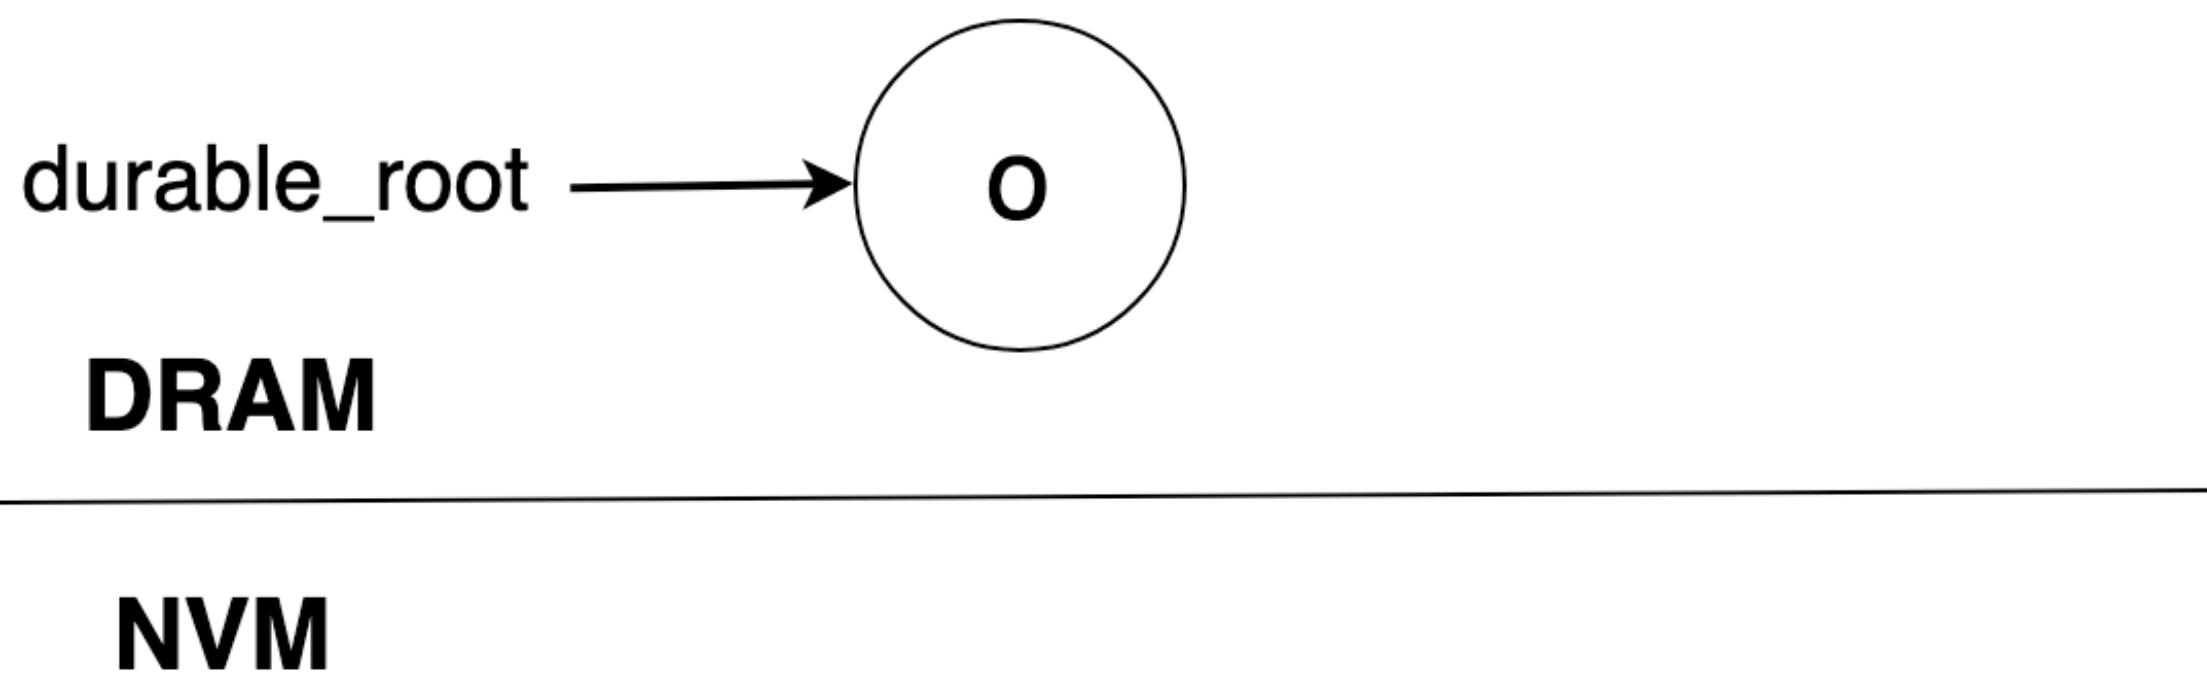
```
do:
```

```
    for x in reachable from o:  
        x.replica = allocate_in_NVM()  
    handshake()  
    for x in reachable from o:  
        copy x to x.replica  
    until all reachable from o are copied
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Concurrent Access

Copier Thread

```
durable_root = o
```

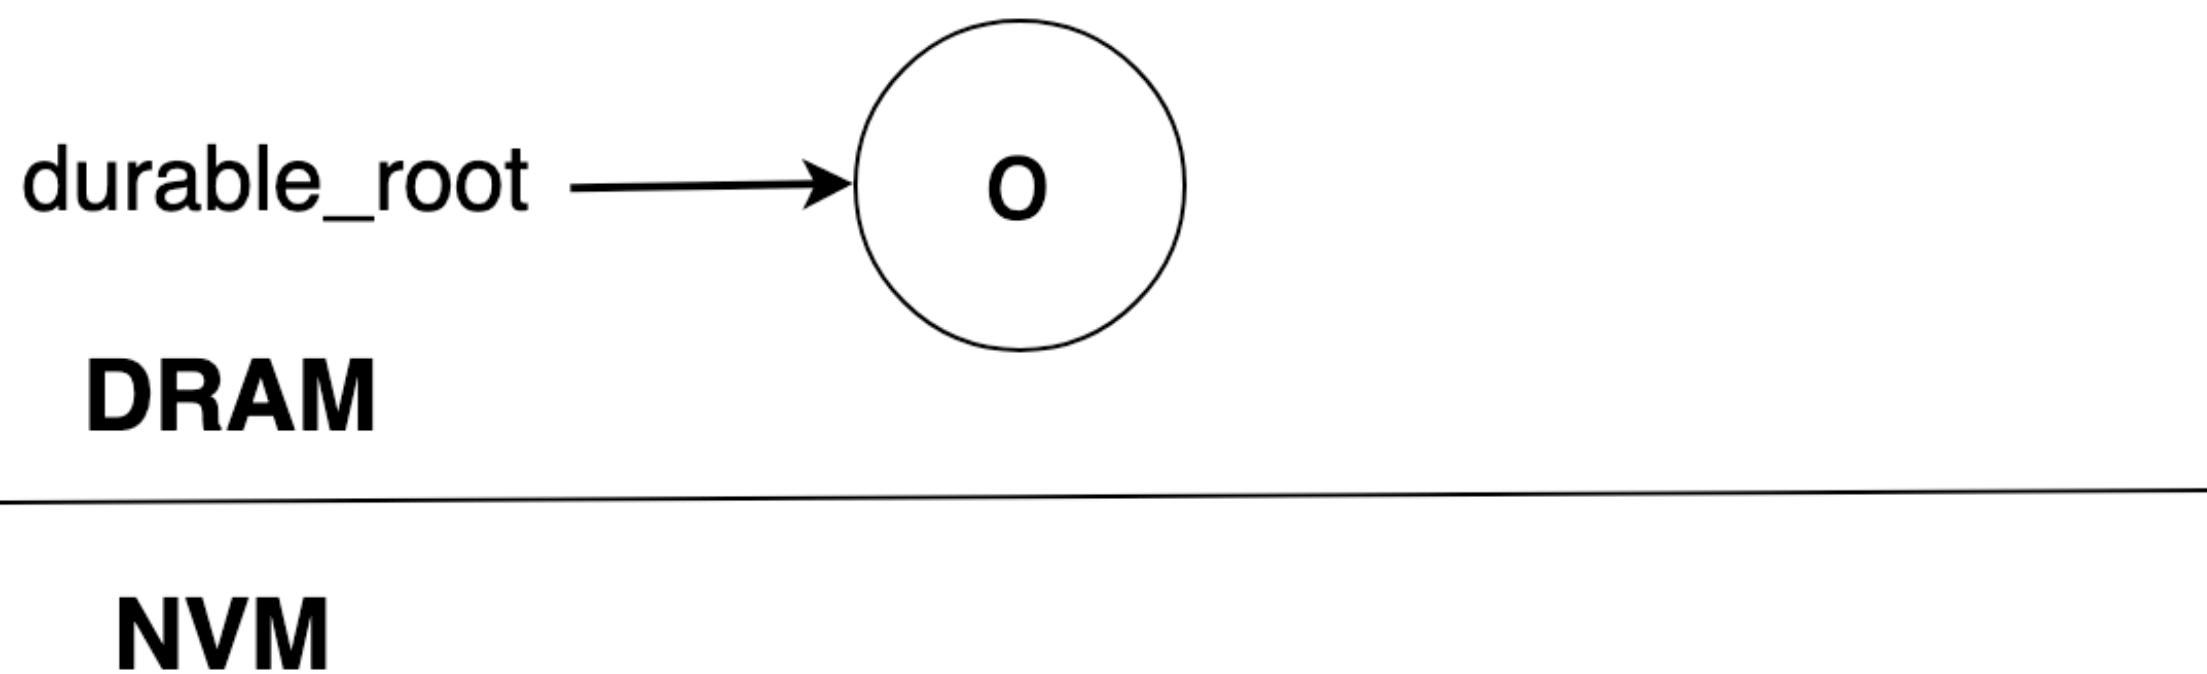
```
do:
```

```
    for x in reachable from o:  
        x.replica = allocate_in_NVM()  
        handshake()  
    for x in reachable from o:  
        copy x to x.replica  
    until all reachable from o are copied
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Concurrent Access

Copier Thread

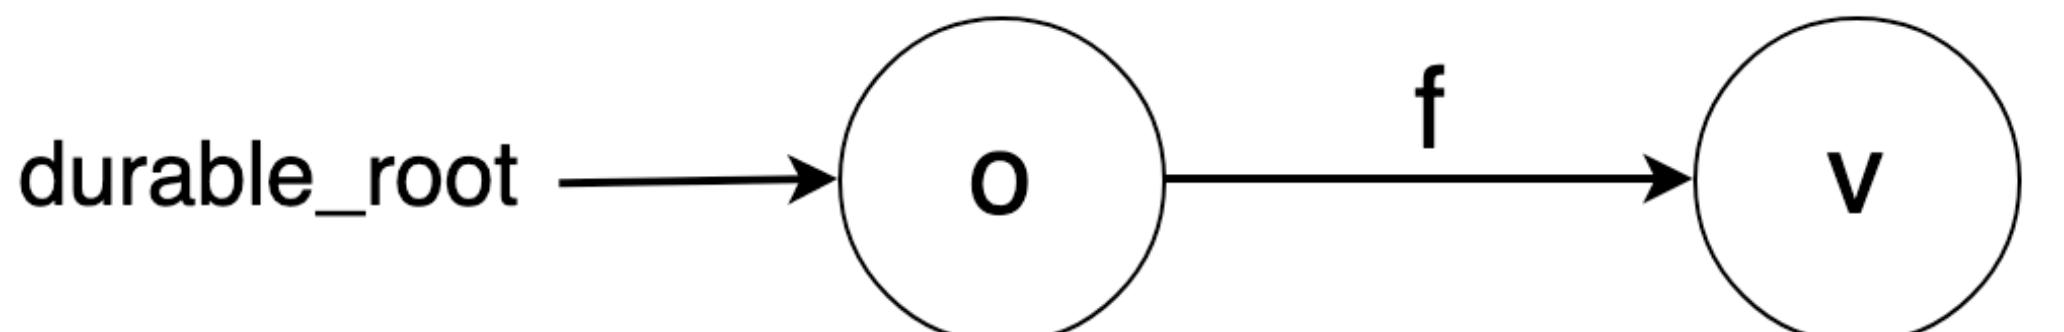
durable_root = o

```
// handshake  
// copy
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

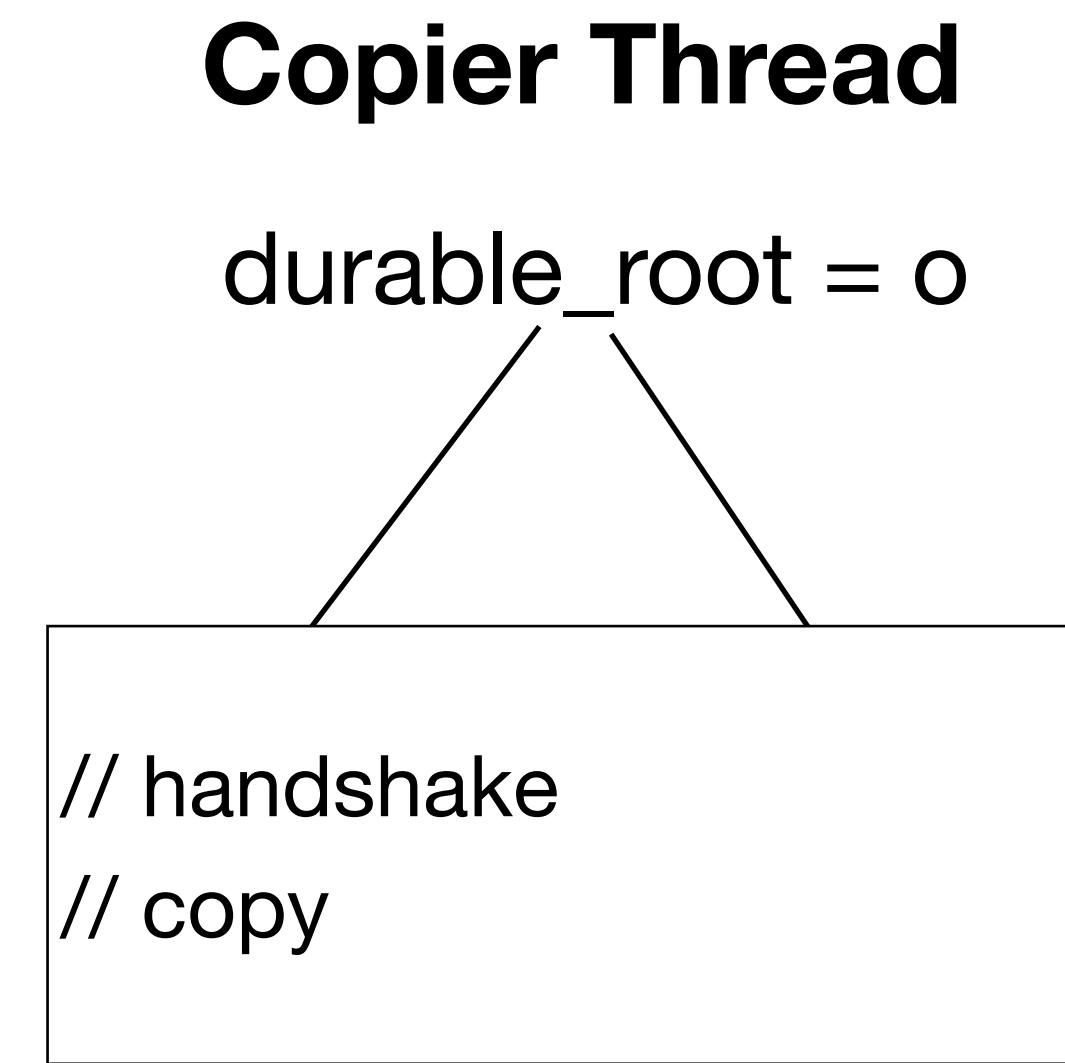
Case 3: Read at Line 4 is reordered with Write at Line 2



DRAM

NVM

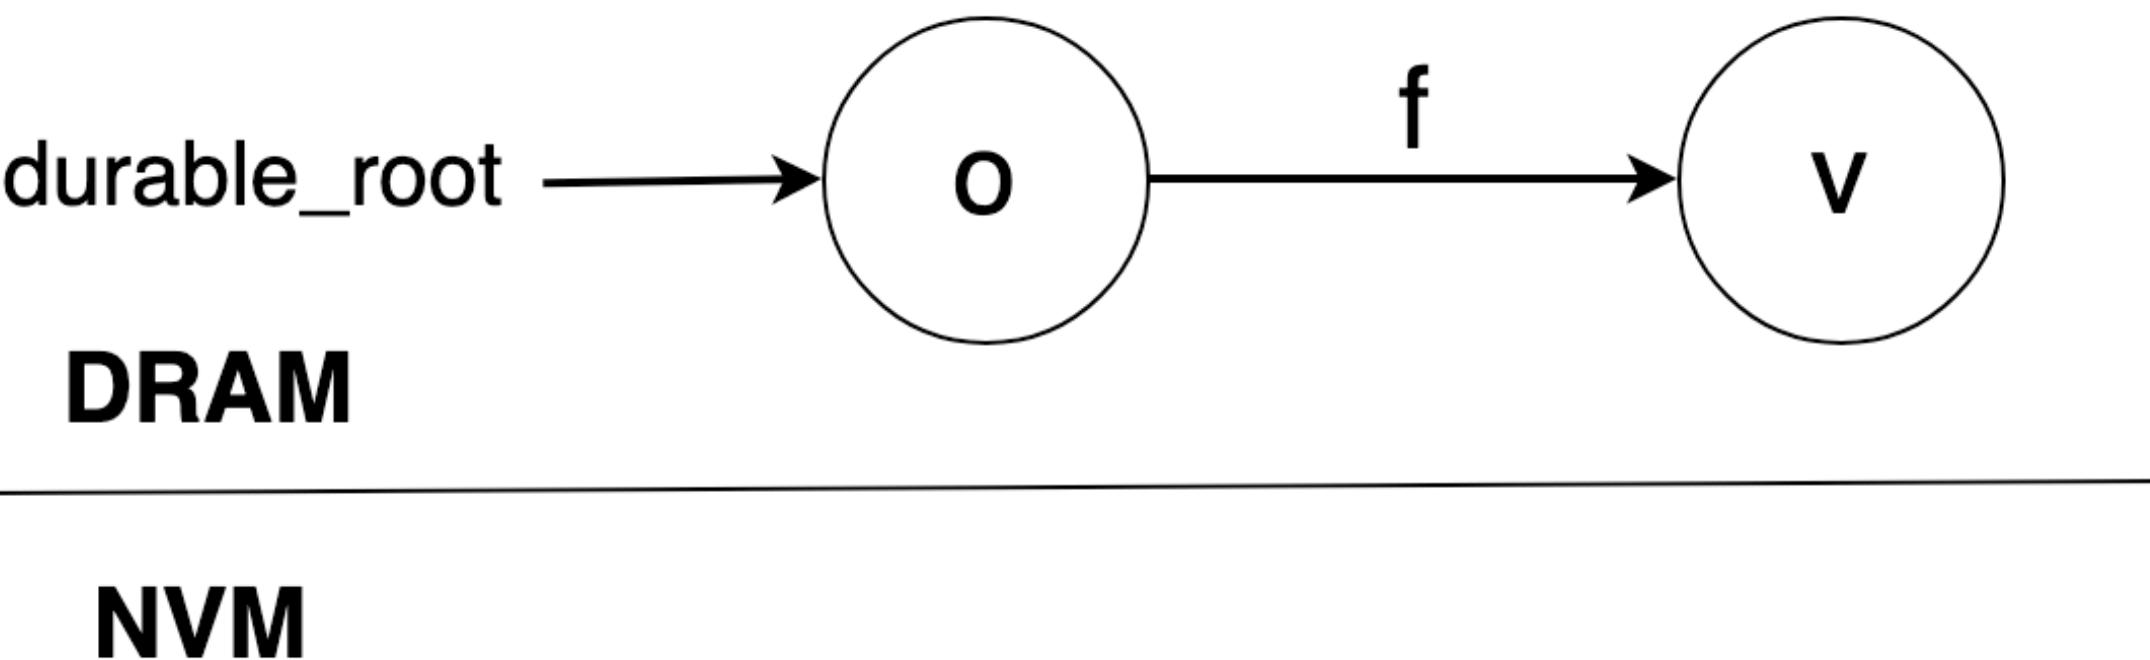
Concurrent Access



Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):
2     o[f] = v
3
4     o' = o.replica
5     if o' == NULL : return
6     make_persistent(v)
7     v' = v.replica
8     o'[f] = v'
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Concurrent Access

Copier Thread

durable_root = o

do:

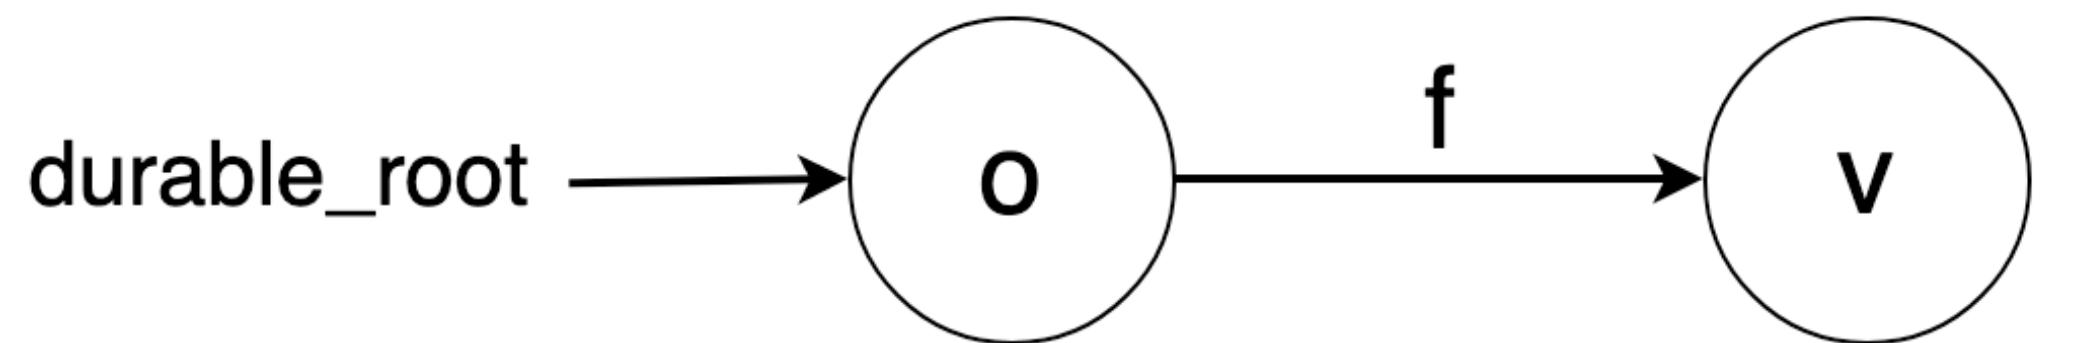
```
for x in reachable from o:  
    x.replica = allocate_in_NVM()  
handshake()  
for x in reachable from o:  
    copy x to x.replica  
until all reachable from o are copied
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

GC safepoint

Case 3: Read at Line 4 is reordered with Write at Line 2



DRAM

NVM

Concurrent Access

Copier Thread

durable_root = o

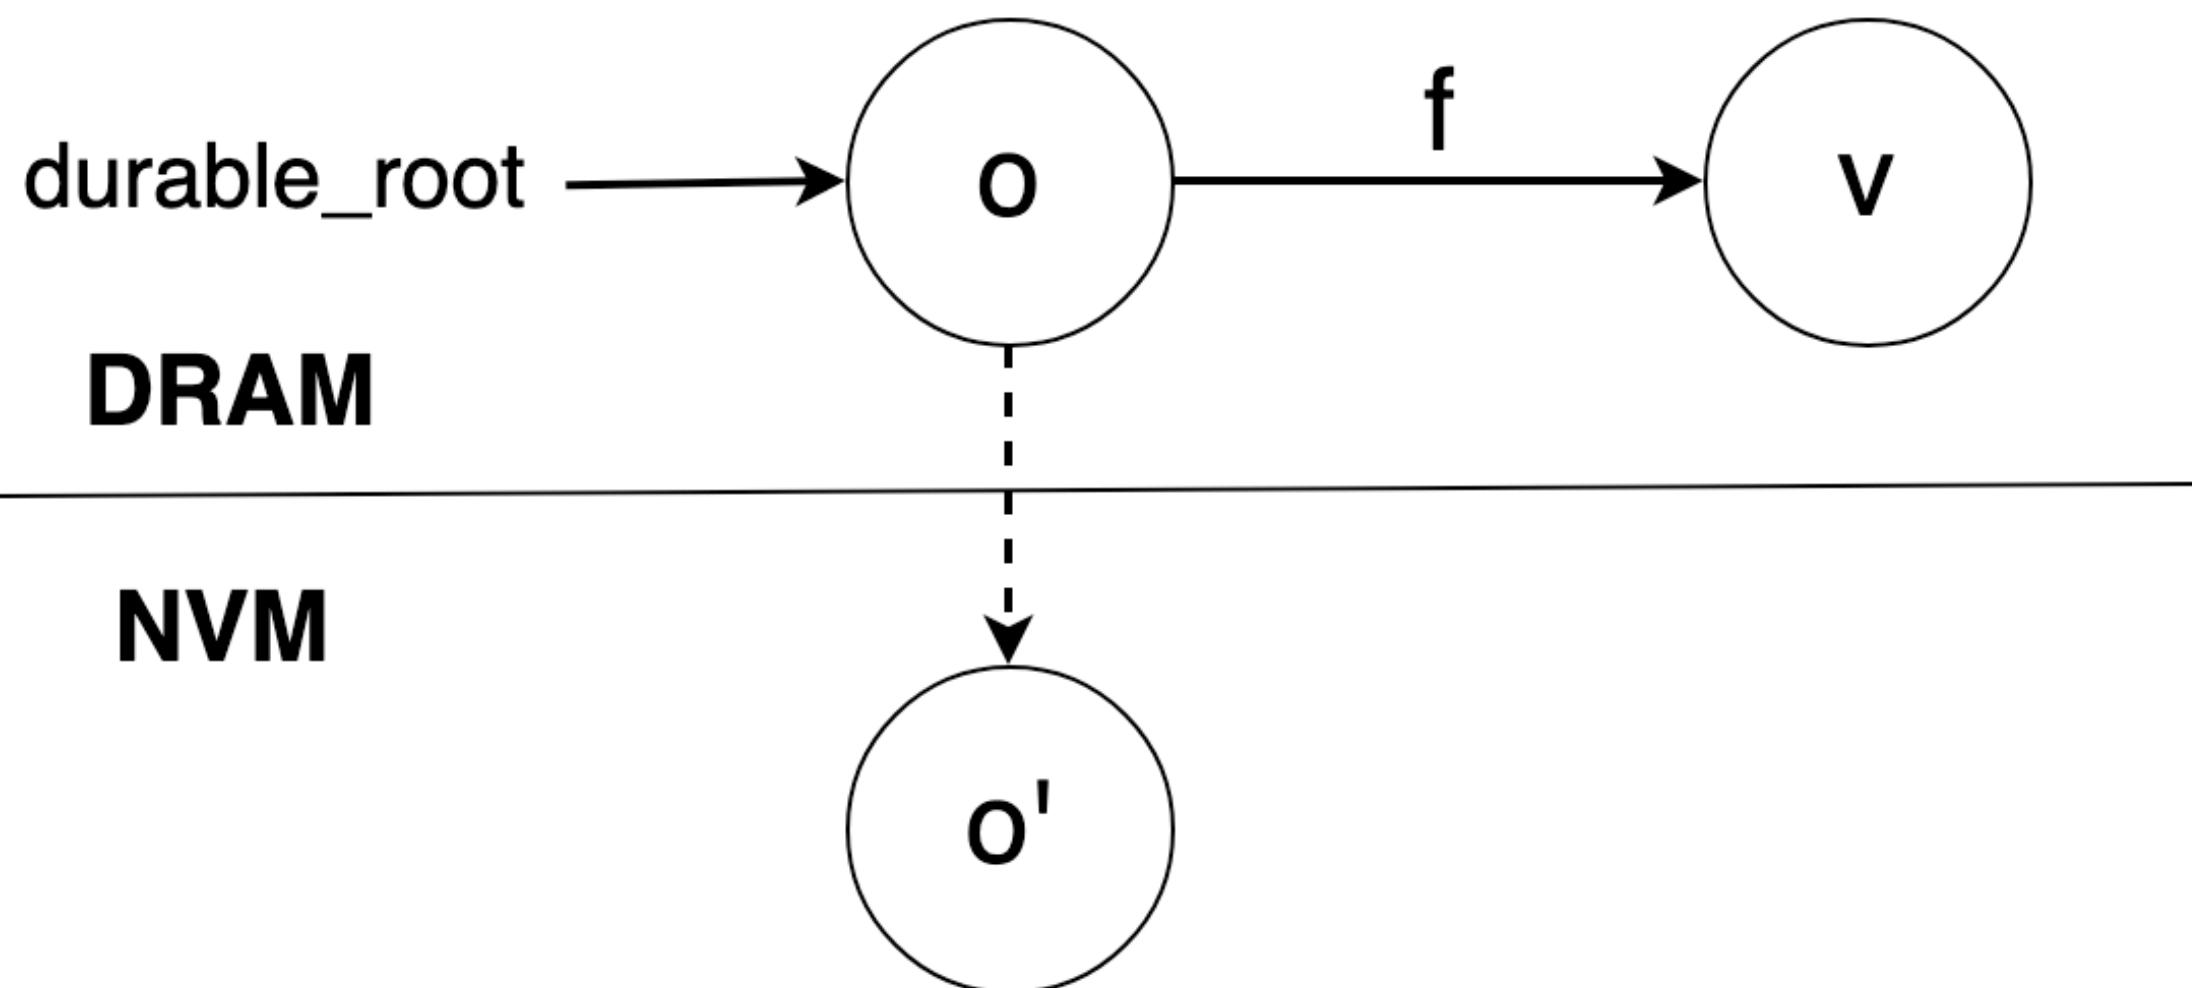
do:

```
for x in reachable from o:  
    x.replica = allocate_in_NVM()  
handshake()  
for x in reachable from o:  
    copy x to x.replica  
until all reachable from o are copied
```

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



Concurrent Access

Copier Thread

durable_root = o

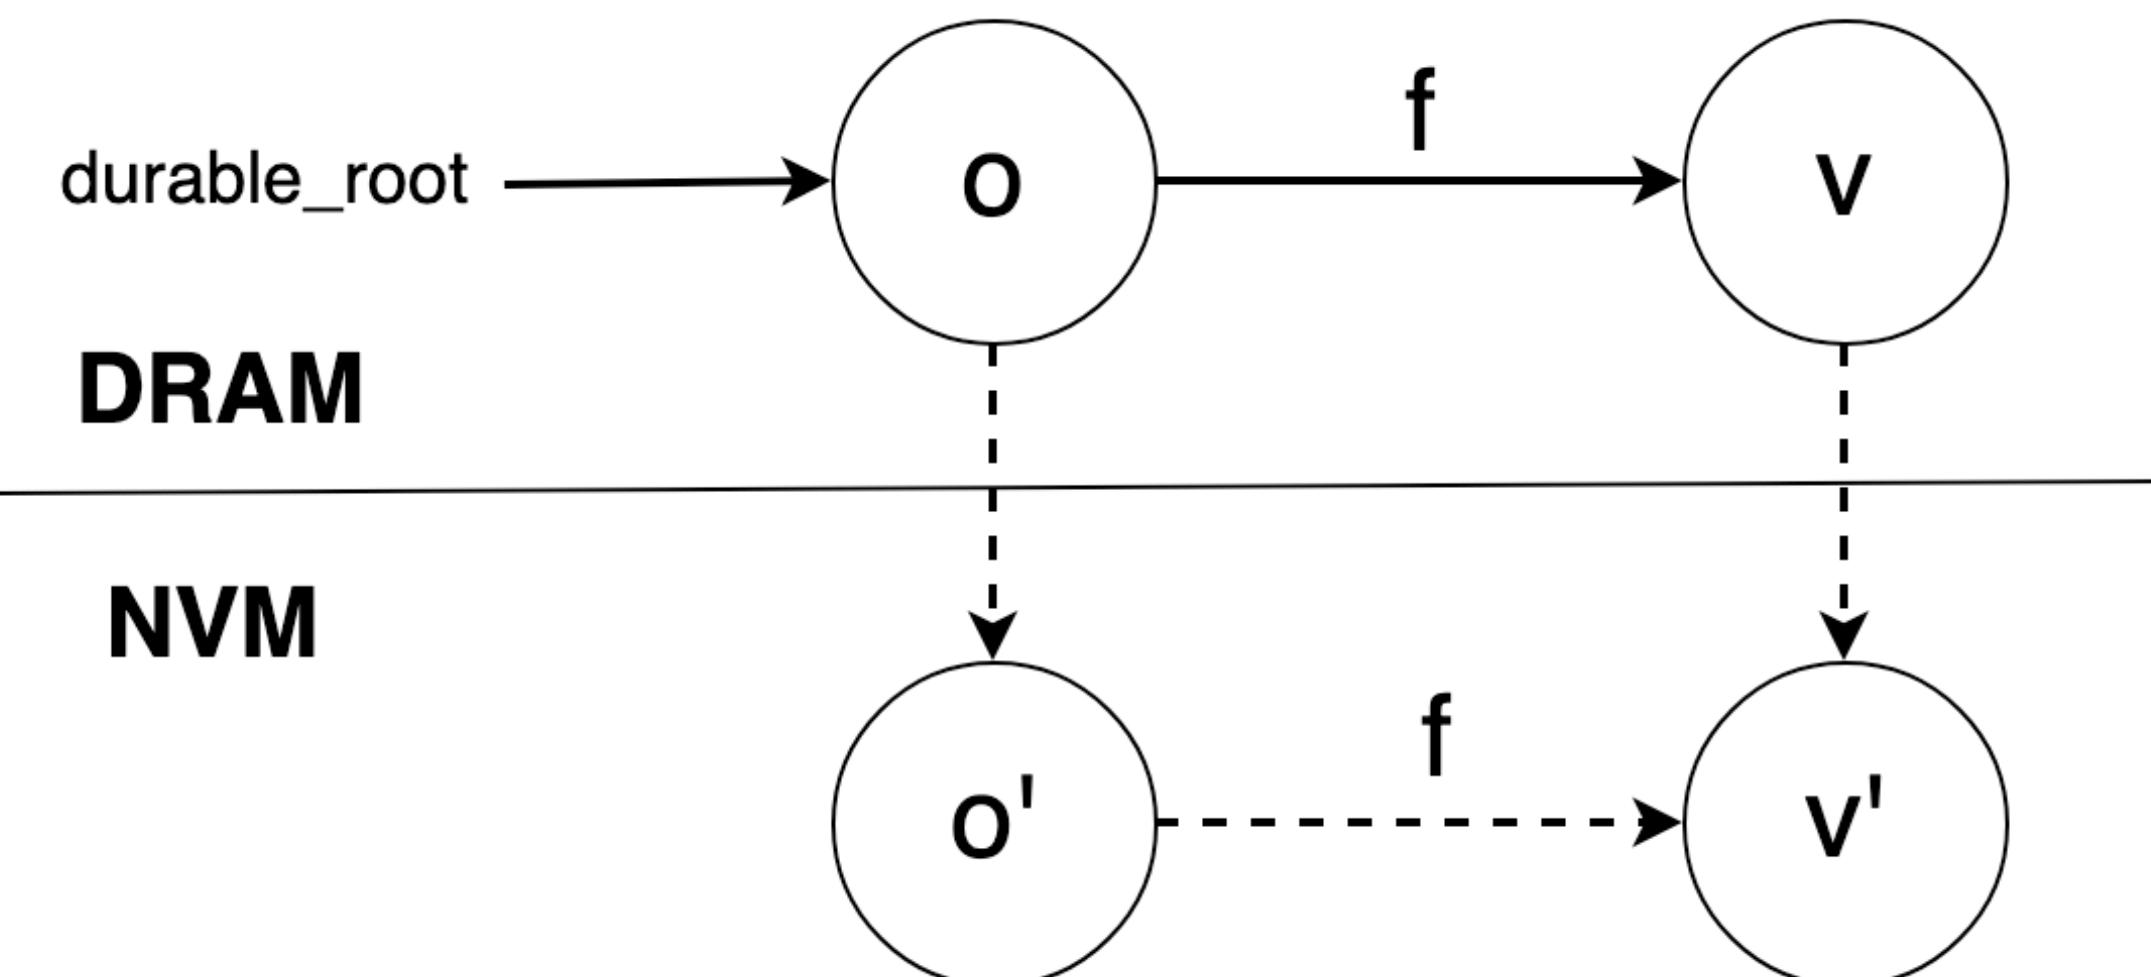
do:

```
for x in reachable from o:  
    x.replica = allocate_in_NVM()  
handshake()  
for x in reachable from o:  
    copy x to x.replica  
until all reachable from o are copied
```

Writer Thread ($o.f = v$)

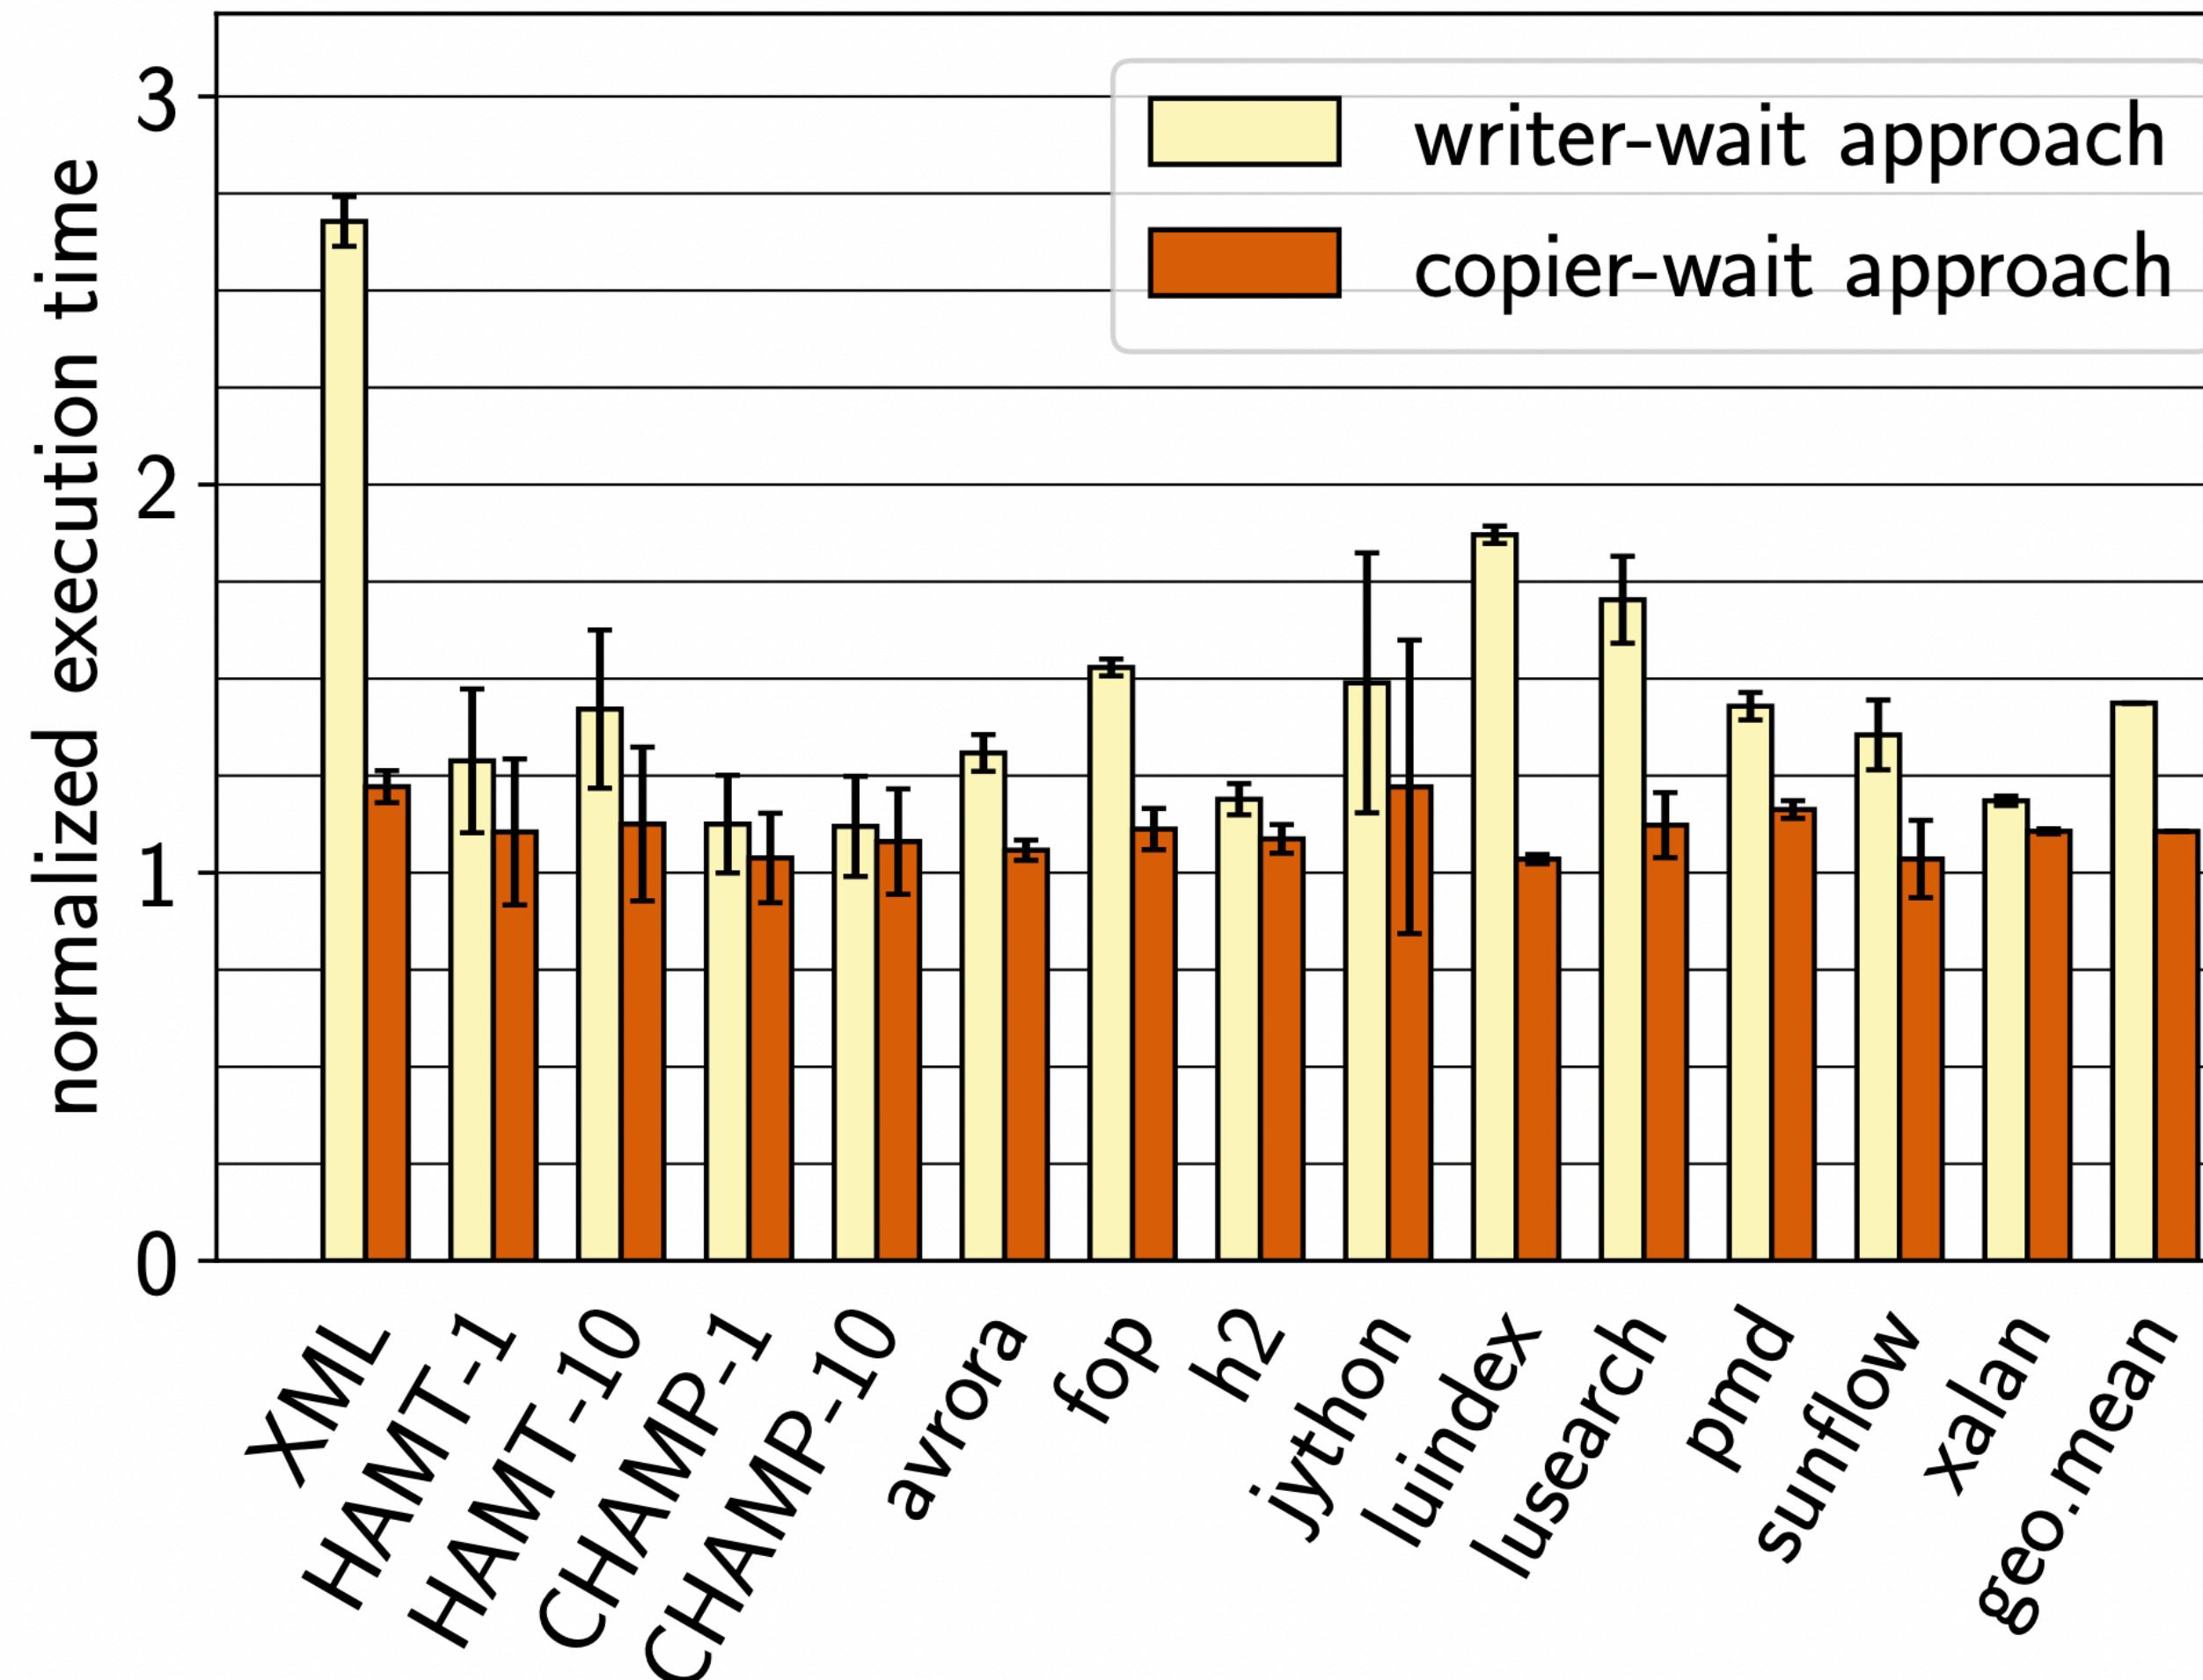
```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

Case 3: Read at Line 4 is reordered with Write at Line 2



**Does shifting overhead to the copier benefit
programs that rarely make objects persistent ?**

No durable roots in program



Elapsed times normalised to standard HotSpot VM

writer-wait approach
43.7% overhead on average

copier-wait approach
10.6% overhead on average

**What about programs that
frequently make objects persistent ?**

Overheads of Copier-wait

- When objects are frequently made persistent, copier-wait approach has high overheads
- Handshake overhead of 37.9 % compared to writer-wait approach when all static fields are annotated as durable roots.

Overheads of Copier-wait

- When objects are frequently made persistent, copier-wait approach has high overheads
- Handshake overhead of 37.9 % compared to writer-wait approach when all static fields are annotated as durable roots.
- Do we always need handshake ?

Do we always need handshake ?

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

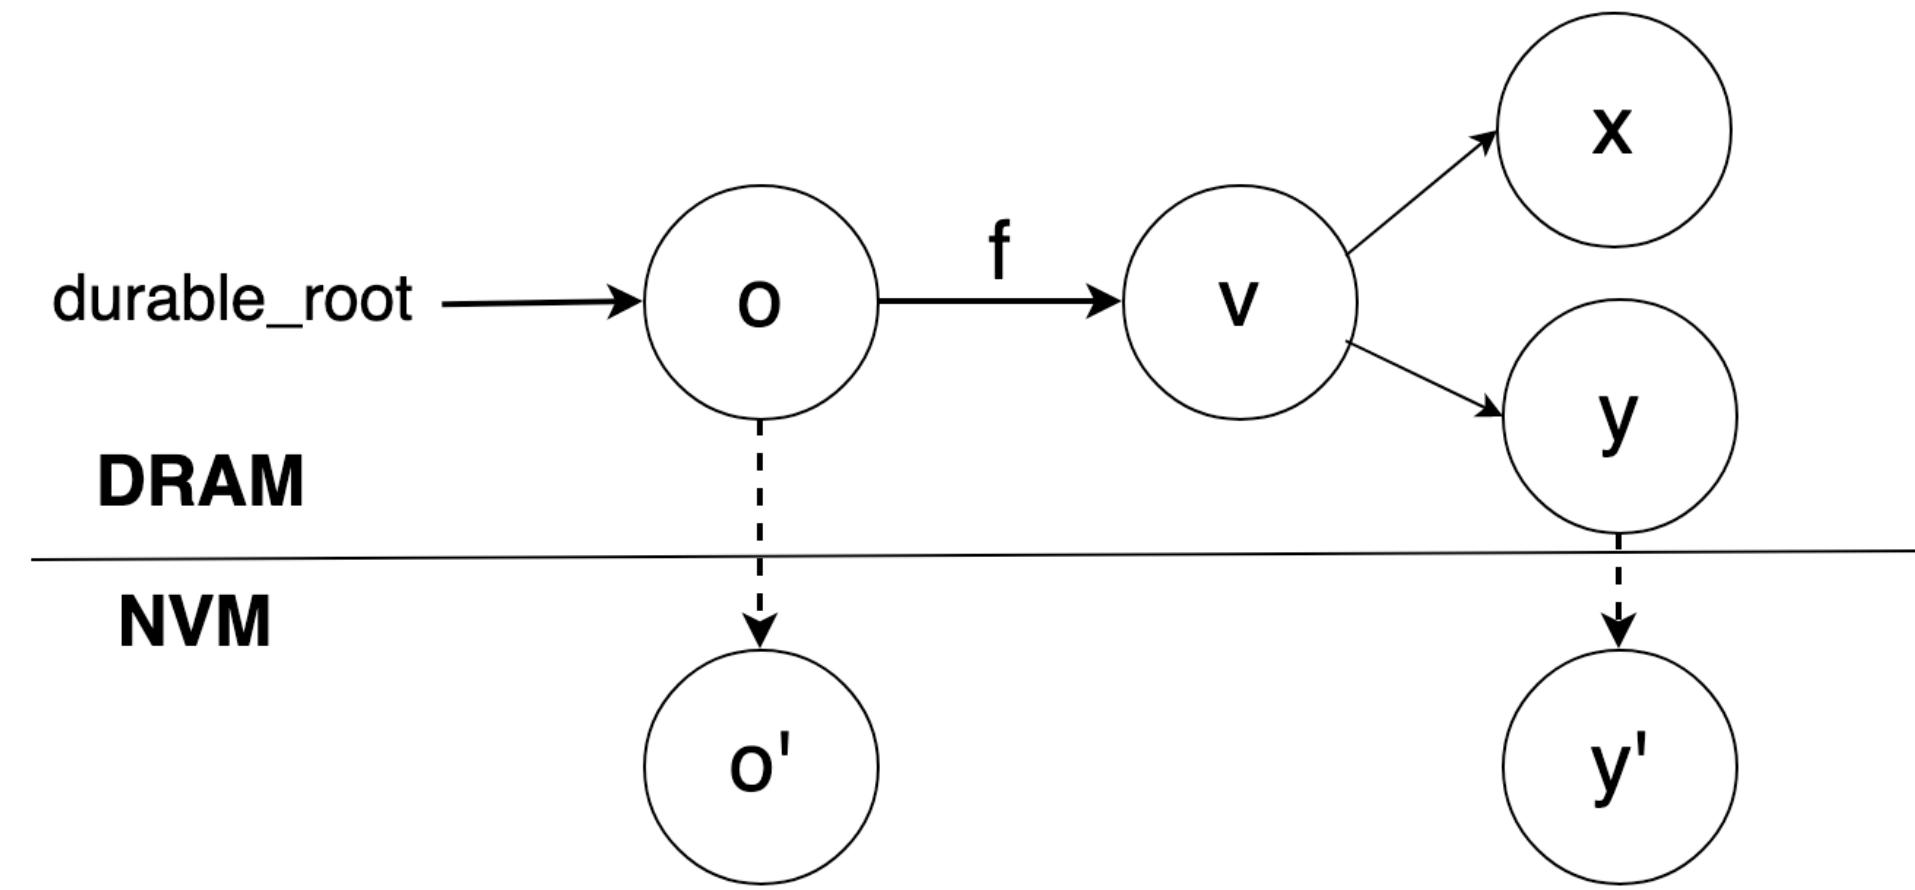
```
do:  
    for x in reachable from o:  
        x.replica = allocate_in_NVM()  
        handshake()  
    for x in reachable from o:  
        copy x to x.replica  
    until all reachable from o are copied
```

Do we always need handshake ?

Writer Thread ($o.f = v$)

```
1 def putfield( o, f, v ):  
2     o[f] = v  
3  
4     o' = o.replica  
5     if o' == NULL : return  
6     make_persistent(v)  
7     v' = v.replica  
8     o'[f] = v'  
9     CLWB(&o'[f])
```

```
do:  
    for x in reachable from o:  
        x.replica = allocate_in_NVM()  
        handshake()  
    for x in reachable from o:  
        copy x to x.replica  
    until all reachable from o are copied
```



If all the objects reachable from v (including) are, either

- Thread Local or
- Persistent

Then handshake can be elided

Persistence-Aware Escape Analysis

- Combined Points-to-Escape analysis is used to identify thread-local abstract objects
- Modify escape analysis to recognise a special abstract object
 - Persistent Object (**P**)

Persistence-Aware Escape Analysis

- Combined Points-to-Escape analysis is used to identify thread-local abstract objects
- Modify escape analysis to recognise a special abstract object
 - Persistent Object (**P**)
- Escape Analysis

`x = A.durable_root`

$x \rightarrow \{ E \}$

`y = x.f`

$y \rightarrow \{ E \}$

`y.f = z`

$y.f \rightarrow \{ E \}$

`t = new Thread(w)`

$t \rightarrow \{ E \} \quad w \rightarrow \{ E \}$

Persistence-Aware Escape Analysis

- Combined Points-to-Escape analysis is used to identify thread-local abstract objects
- Modify escape analysis to recognise a special abstract object
 - Persistent Object (**P**)
- Persistence-Aware Escape Analysis

`x = A.durable_root`

$x \rightarrow \{ P \}$

`y = x.f`

$y \rightarrow \{ P \}$

`y.f = z`

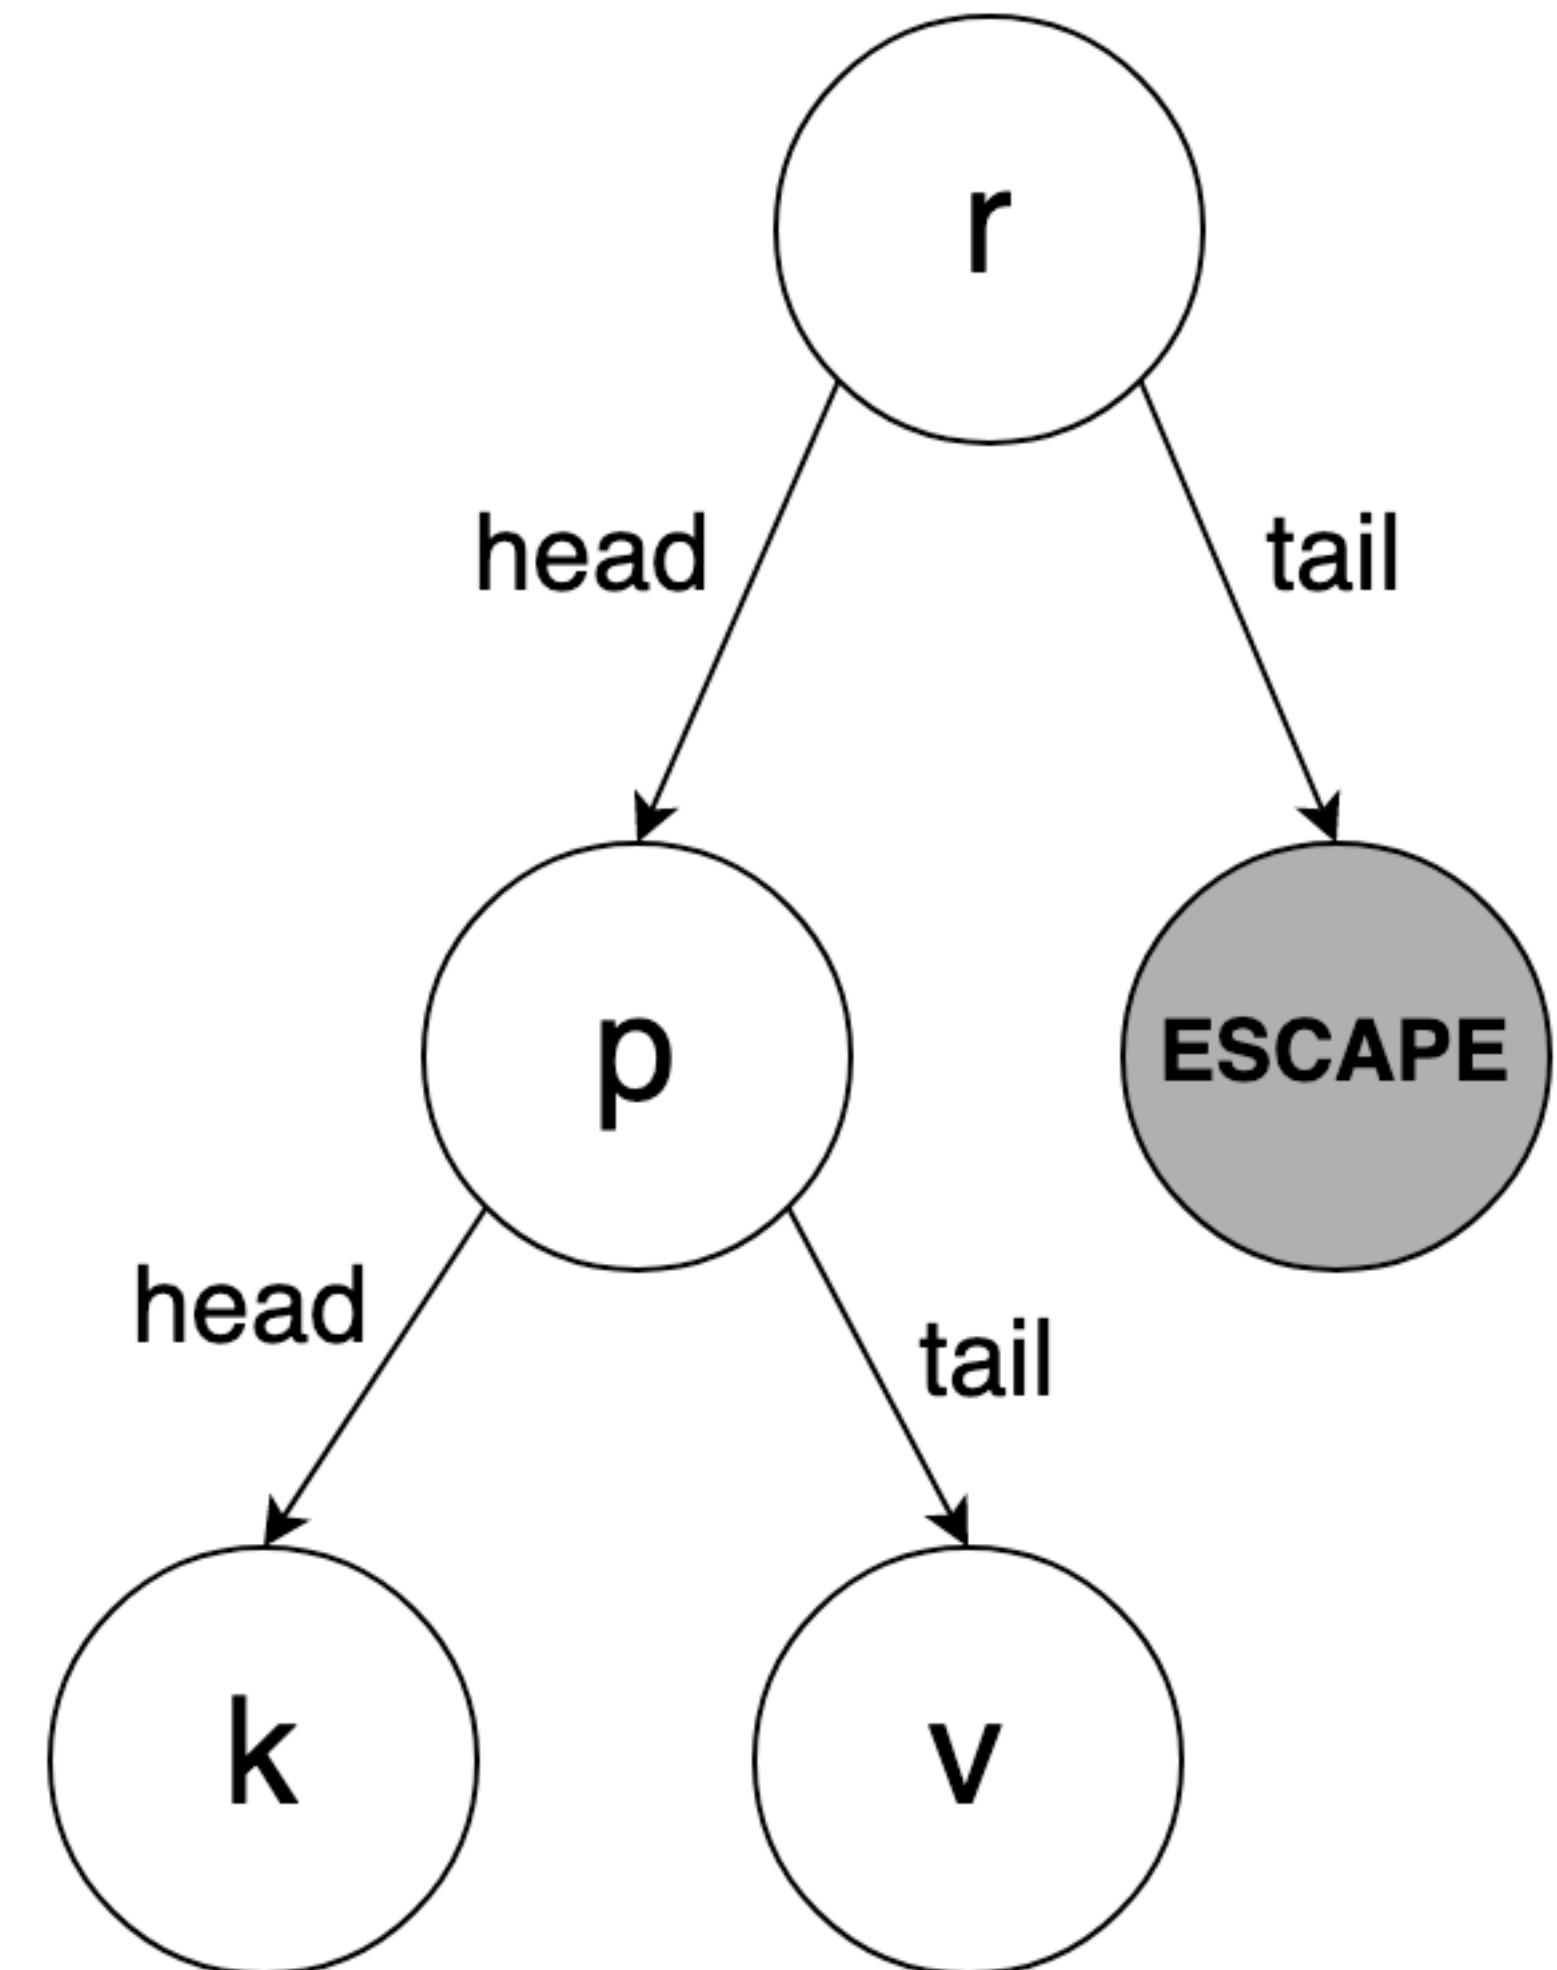
$y.f \rightarrow \{ P \}$

`t = new Thread(w)`

$t \rightarrow \{ E \} \quad w \rightarrow \{ E \}$

Example

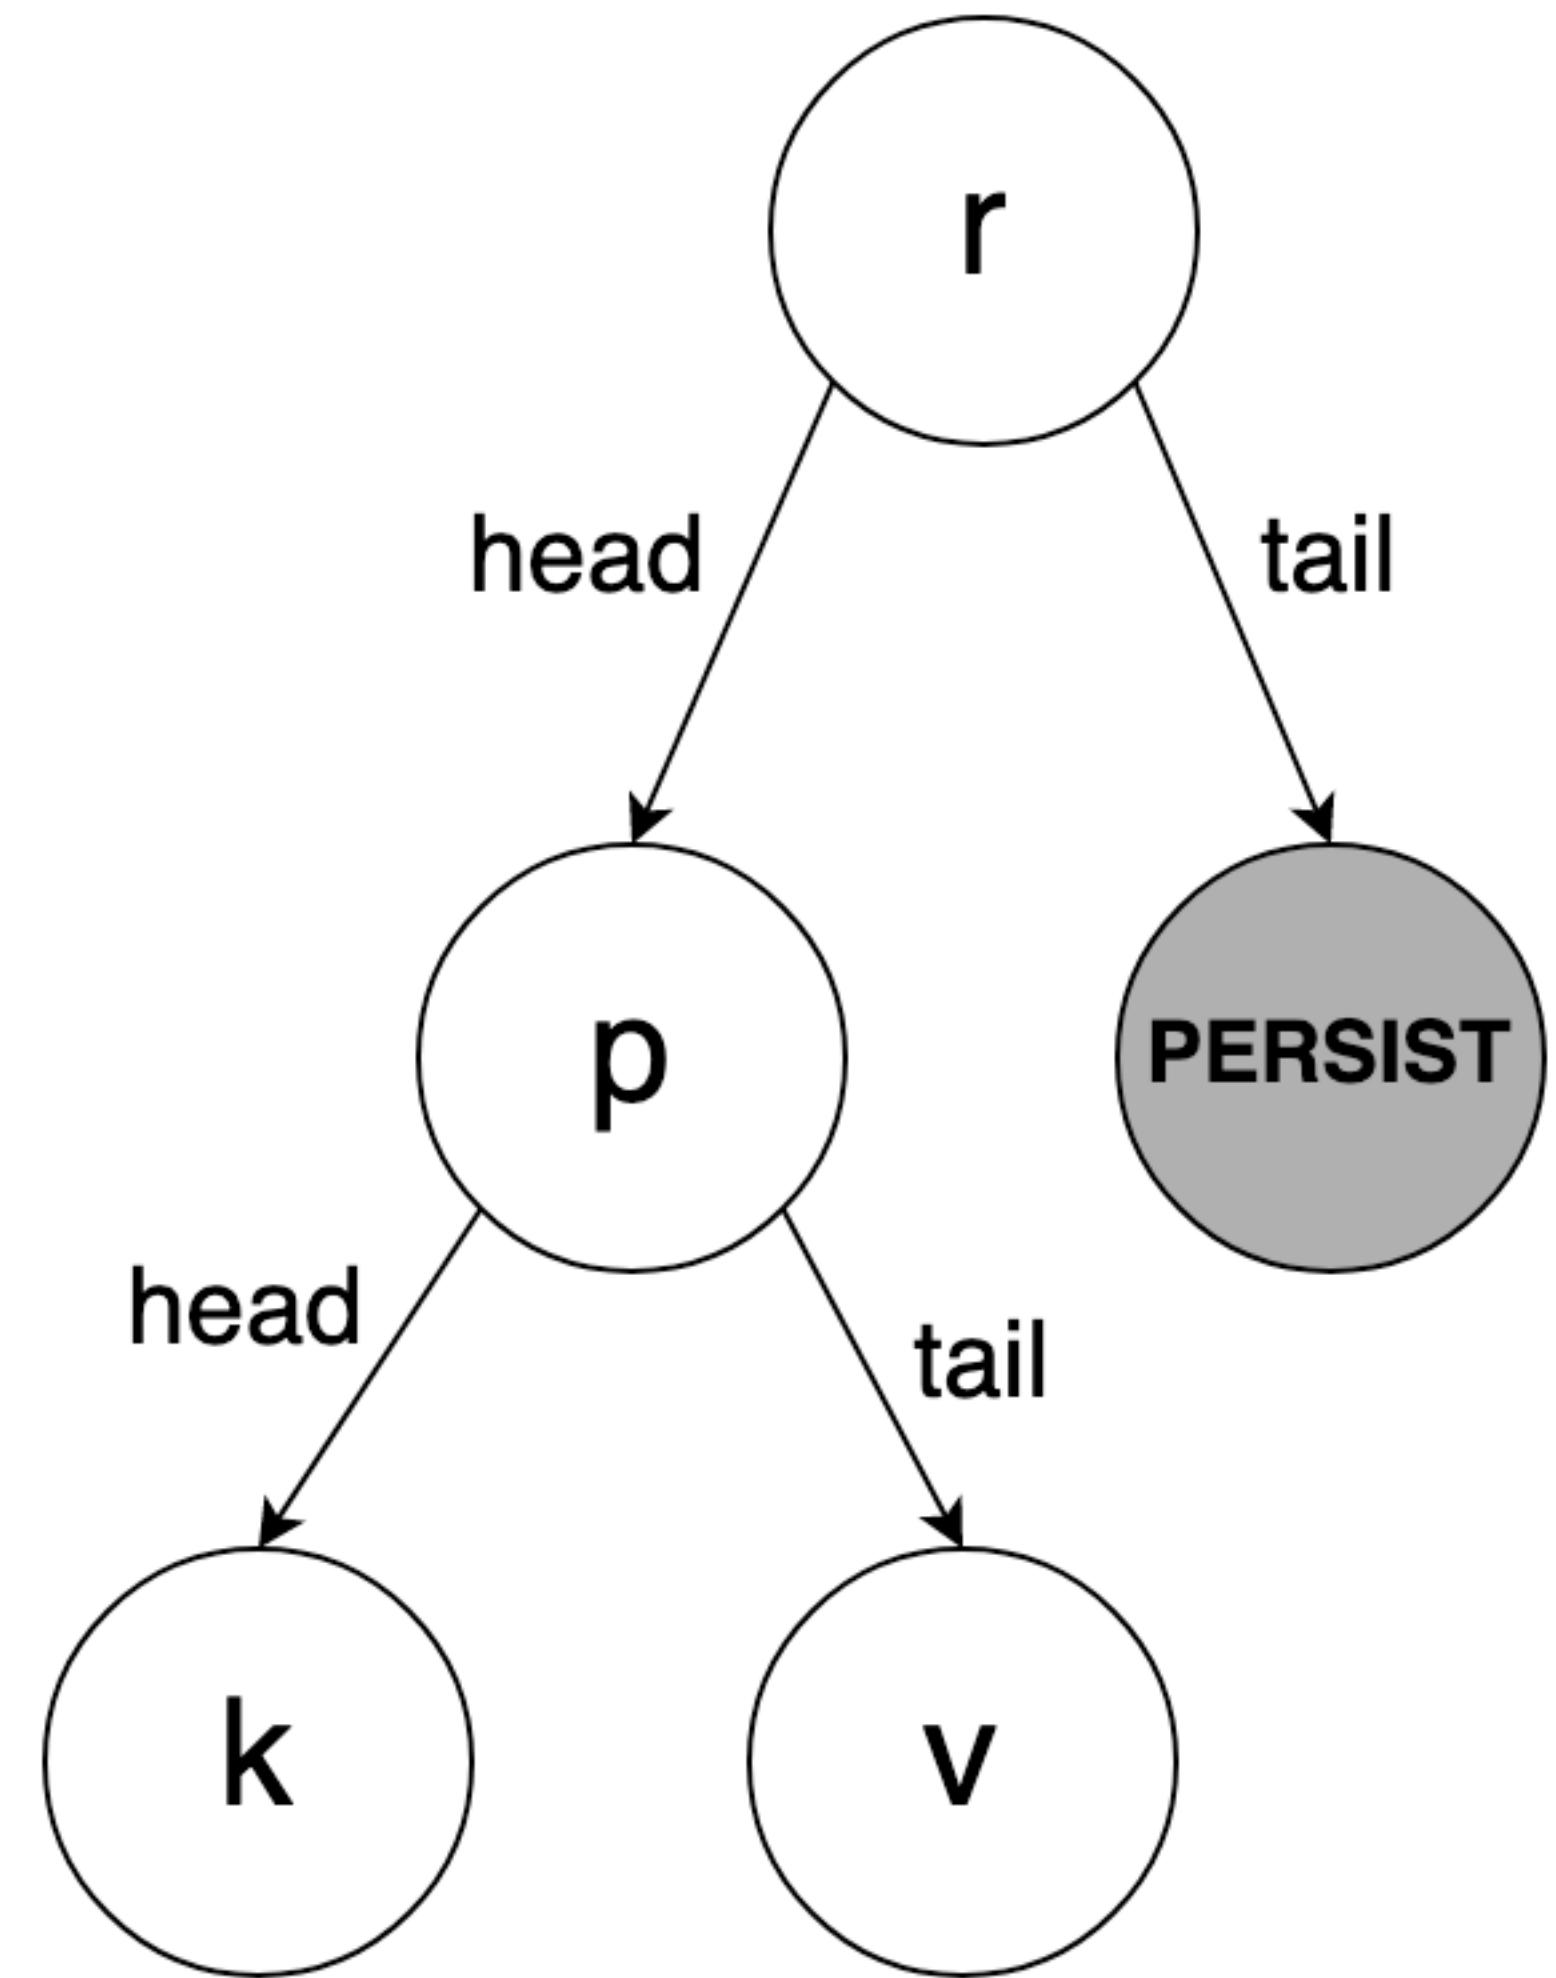
```
1 class Program {  
2     @durable_root static Pair assocList;  
3     @durable_root static int count;  
4     static void assoc(Object k, Object v) {  
5         Pair p = new Pair(); // p is volatile object  
6         p.head = k;  
7         p.tail = v;  
8         Pair r = new Pair(); // r is volatile object  
9         r.head = p;  
10        r.tail = Program.assocList;  
11        Program.assocList = r; // p and r become persistent  
12        Program.count = Program.count + 1; } }
```



White objects are thread-local objects

Example

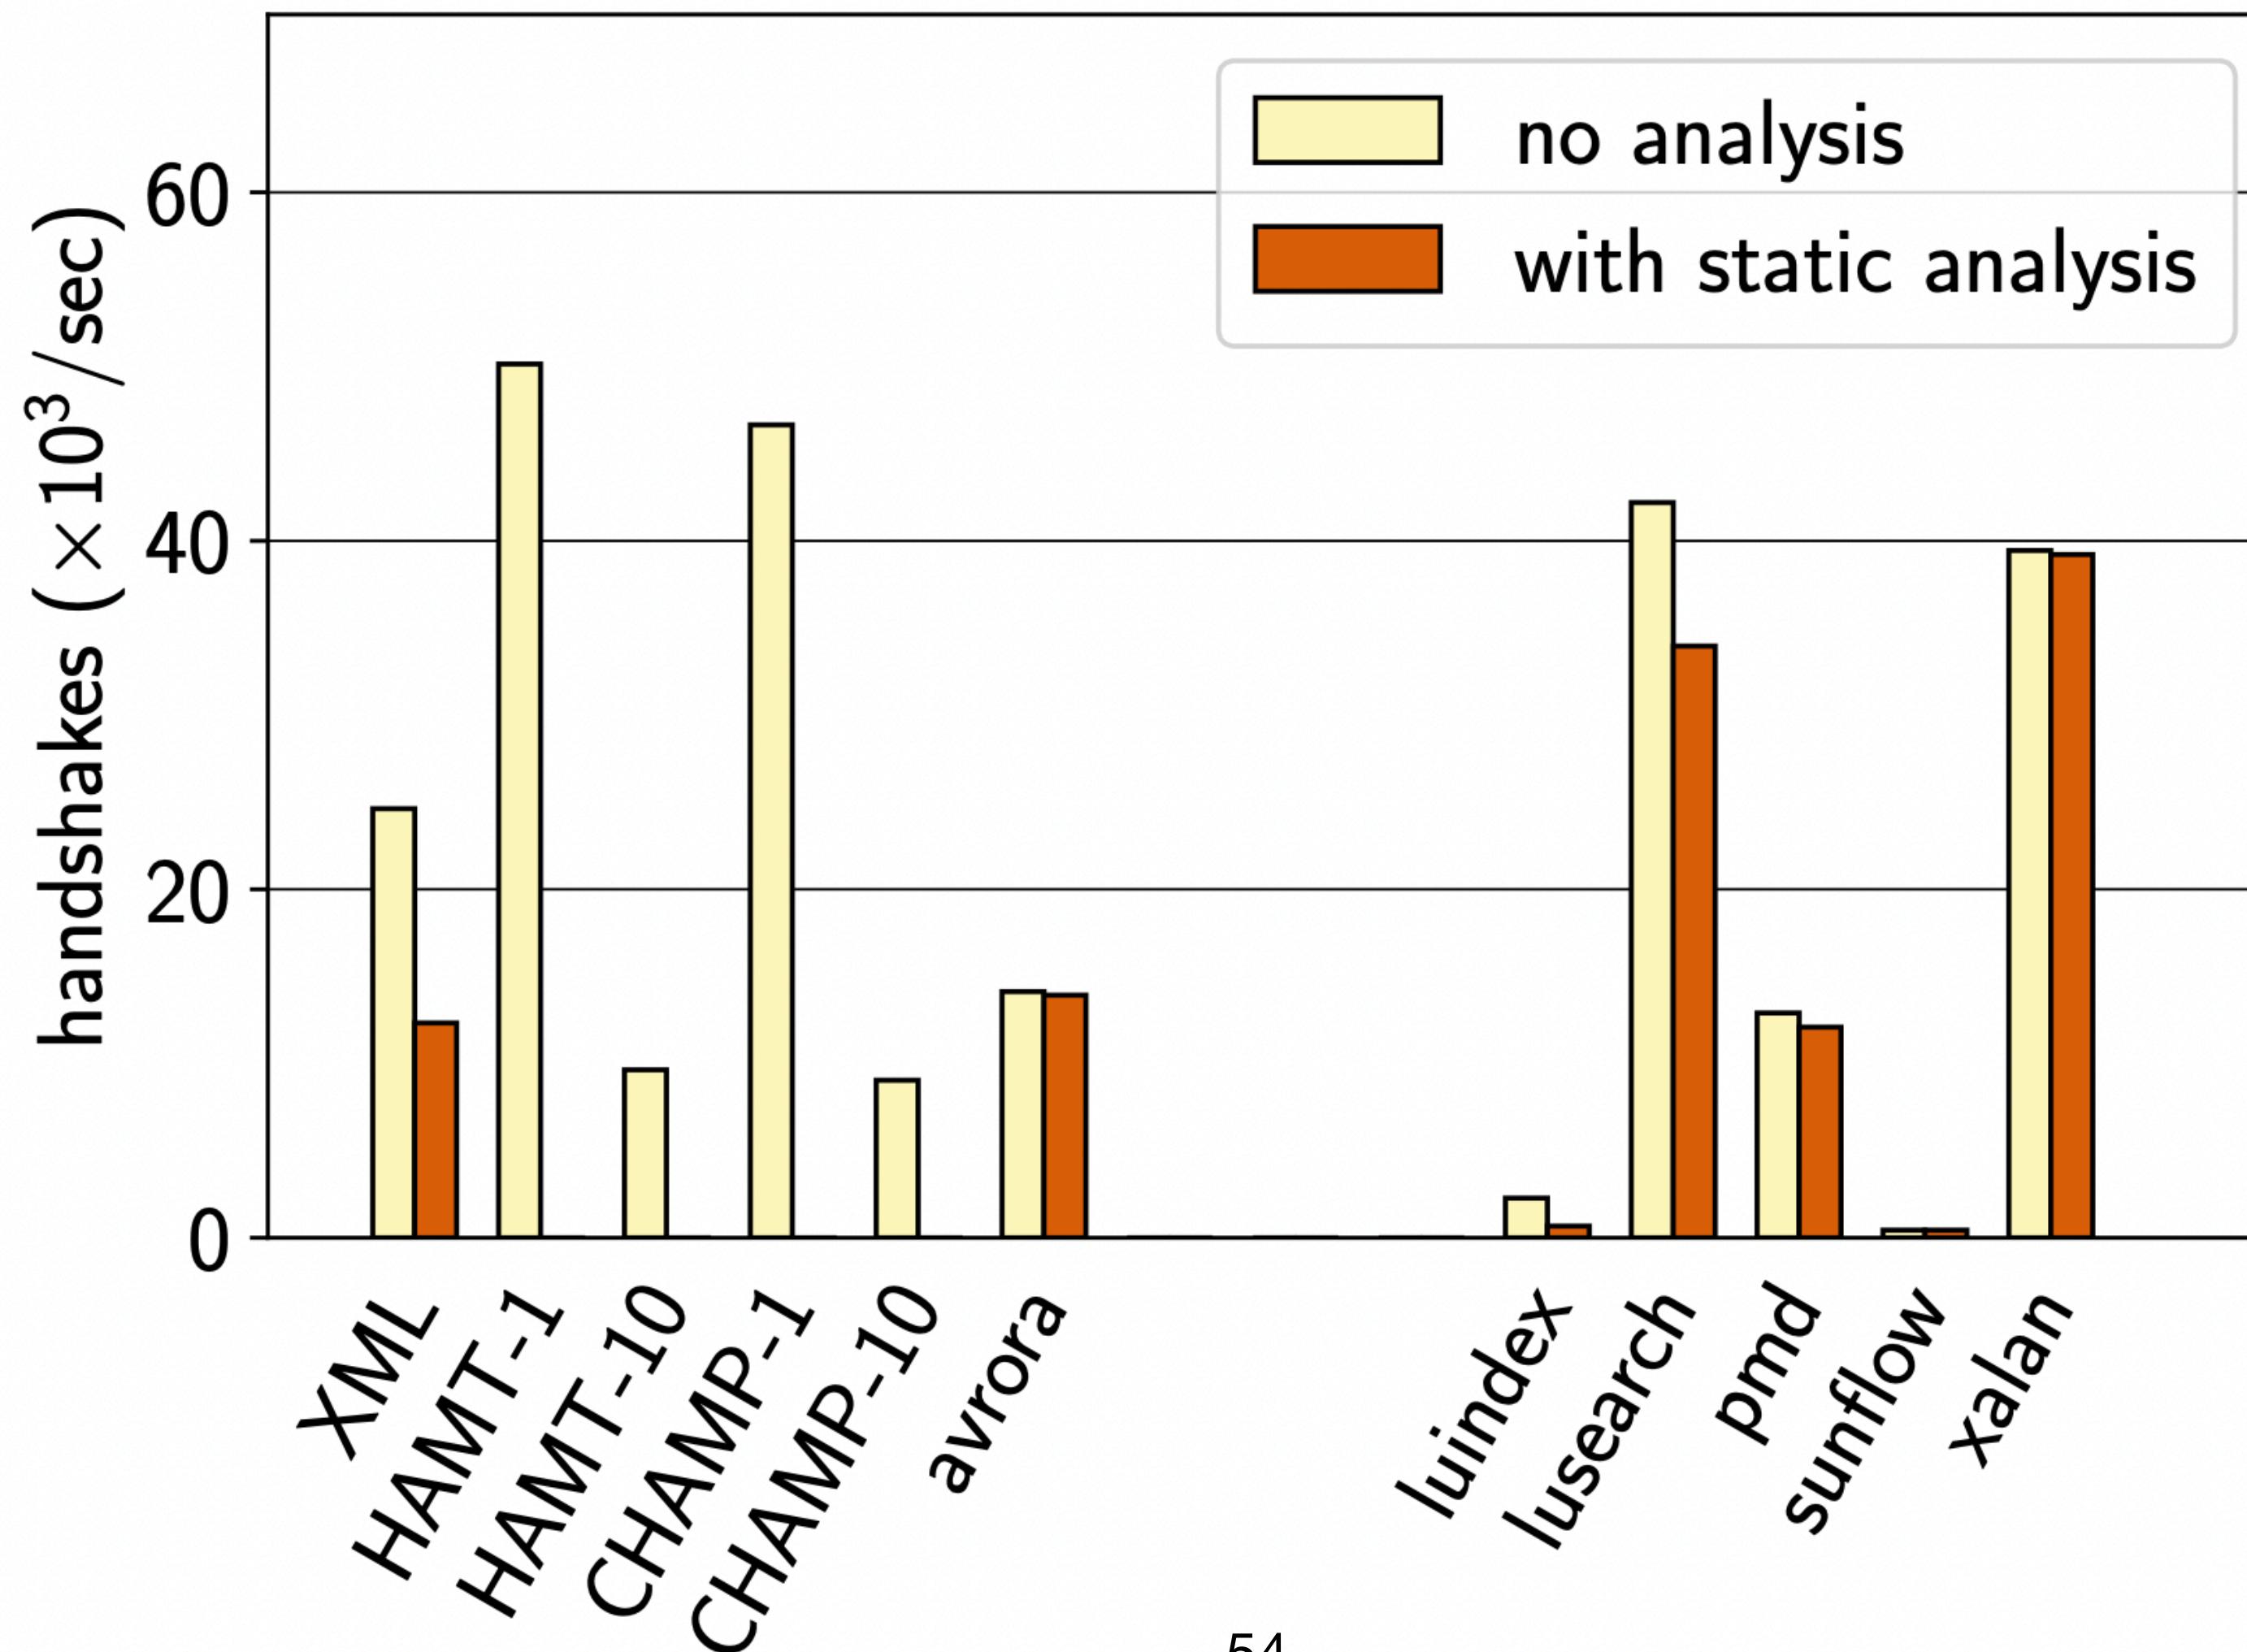
```
1 class Program {  
2     @durable_root static Pair assocList;  
3     @durable_root static int count;  
4     static void assoc(Object k, Object v) {  
5         Pair p = new Pair(); // p is volatile object  
6         p.head = k;  
7         p.tail = v;  
8         Pair r = new Pair(); // r is volatile object  
9         r.head = p;  
10        r.tail = Program.assocList;  
11        Program.assocList = r; // p and r become persistent  
12        Program.count = Program.count + 1; } }
```



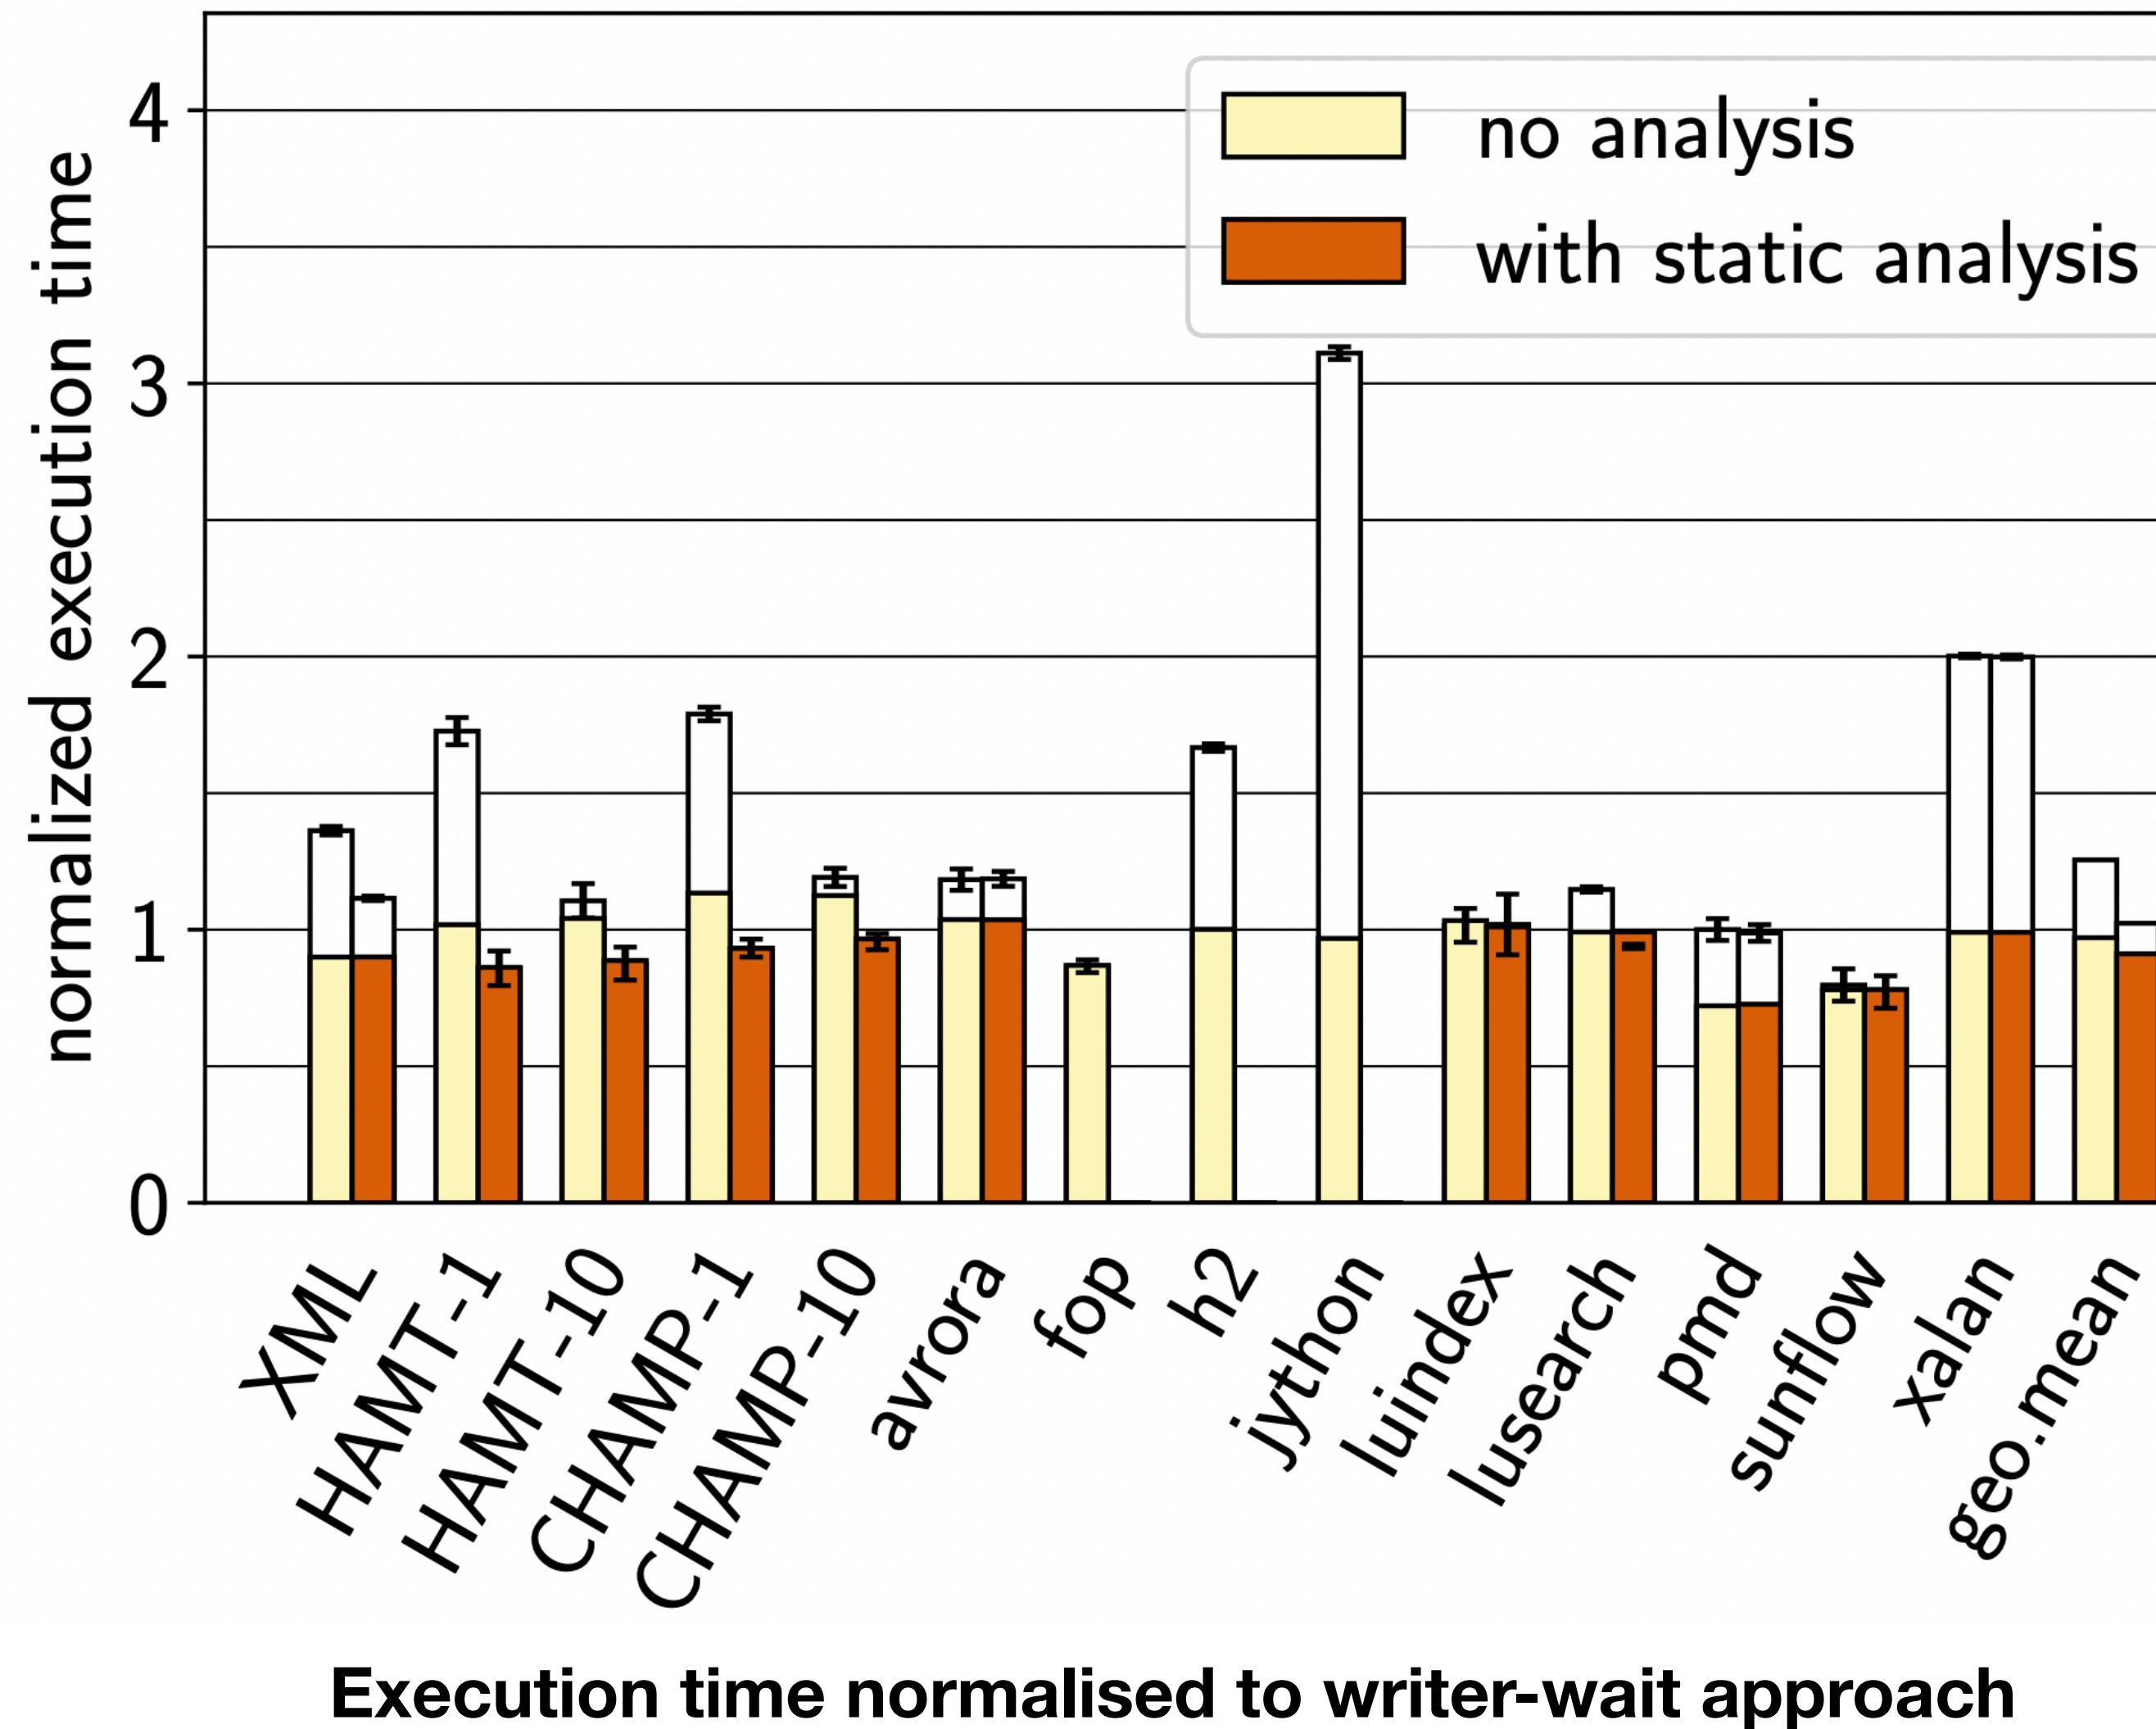
White objects are thread-local objects

**Was static analysis successful in
eliminating handshakes ?**

Number of handshakes per second



Execution time



Static analysis eliminated 52% overheads on average

Copier-wait slower by 2.4% on average

Related Work

- AutoPersist [Shull et al. 2019] uses copier flag that is accessed atomically. Copier fails when race is detected
- QuickCheck [Shull et al. 2019] and P-INSPECT [Kokolis et al. 2020] try to reduce the write barrier overheads but fail to handle races.
- StaticPersist [Bansal 2023] uses static analysis to verify if programmer has correctly made the objects persistent (No pointer from persistent object to volatile object)

Covered in the paper

- Explanation of how the transitive closure of an object is copied.
- Handle race between copier and writer [Ragged Synchronization]
- Synchronization for multiple copiers trying to copy same object
- Correctness Argument
- Additional overheads eliminated by Persistence-aware Escape analysis
- Flow functions for Persistence-aware Escape analysis

Conclusion

- Shifted the overheads from writer thread to copier thread
- 23% performance improvement on average for programs that rarely make objects persistent
- Static analysis decreases 52% of persistence-related overheads in the copier-wait approach
- Performance of copier-wait approach comparable (better in some case) to writer-wait approach

Conclusion

- Shifted the overheads from writer thread to copier thread
- 23% performance improvement on average for programs that rarely make objects persistent
- Static analysis decreases 52% of persistence-related overheads in the copier-wait approach
- Performance of copier-wait approach comparable (better in some case) to writer-wait approach

Thank You